

End-to-End Data Science using Watson Studio

April 30, 2020

The session starts at 9:00am.

End-to-End Data Science using IBM's Watson Studio




Power of data. Simplicity of design. Speed of innovation.

Bernie Beekman
Michael Cronk
James Parry

Agenda




Time	Description
9:00 AM – 10:00 AM	Overview of Watson Studio Lab Orientation 1,2
10:00 AM – 10:15 AM	Break
10:15 AM – 11:45 PM	Lab 1-2 – Set up Environment, Watson Knowledge Catalog
11:45 PM – 12:45 PM	Lab Review 1,2 /Lab Orientation 3,4 Lunch
12:45 PM – 02:00 PM	Lab 3,4 – Data Refinery, SPSS Modeler
02:00 PM – 02:30 PM	Lab Review 3,4 / Lab Orientation 5,6
02:30 PM – 03:30 PM	Lab 5,6 – Machine Learning with SparkML, AutoAI+DevOps
03:30 PM – 04: 00 PM	Lab Review 5,6 / Lab Orientation 7
04:00 PM – 05:00 PM	Lab 7 – Watson OpenScale

The Zoom Tool Bar




We will stay on mute during the main sessions. You may unmute yourself during breakouts.

We will not use video to preserve bandwidth during our class today.

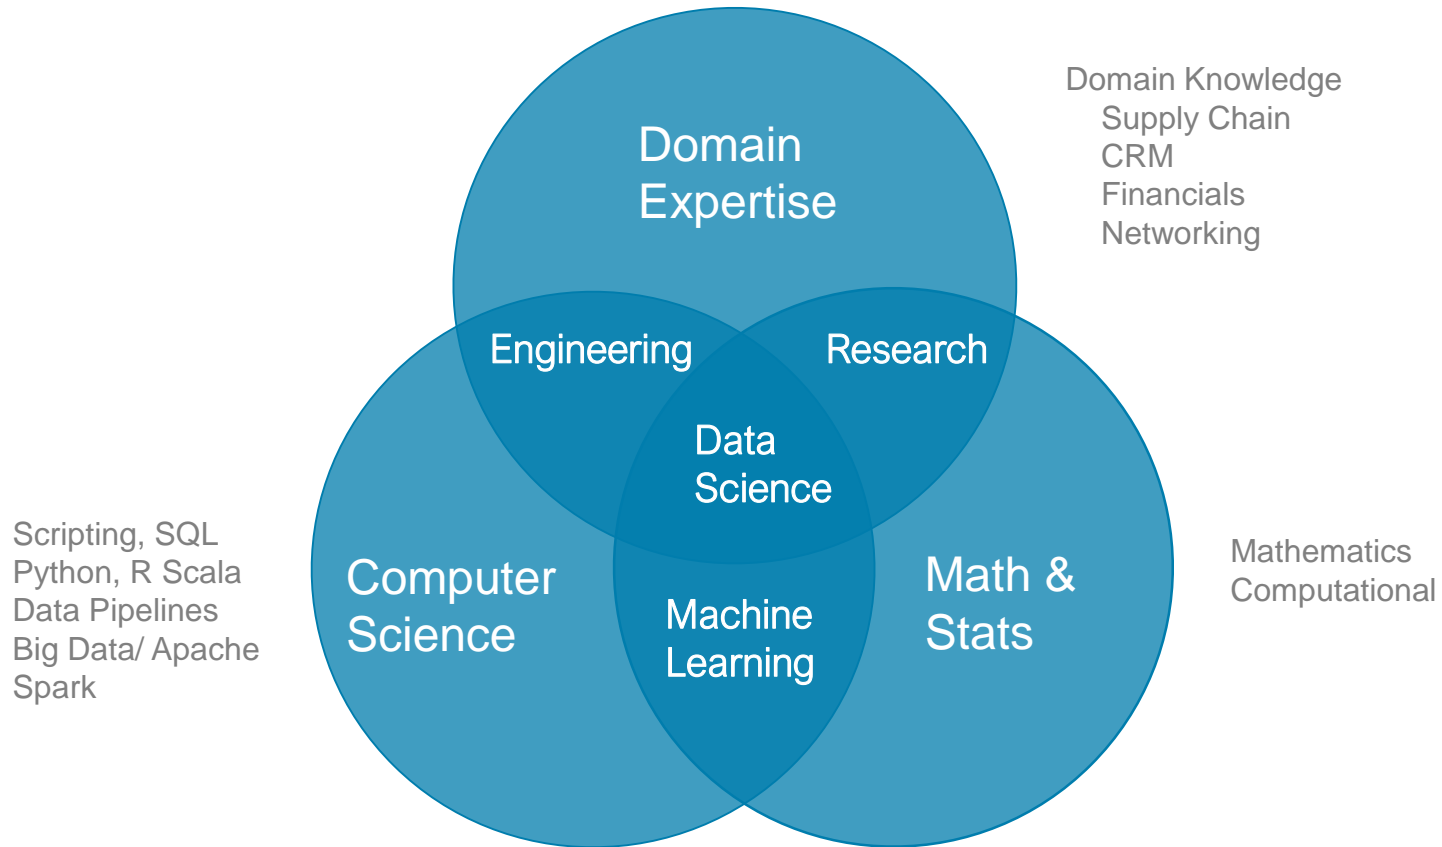


The Share Screen function has been deactivated during the lectures

Outline

- **Data Science Overview** 
- **Watson Studio Overview**
- **Lab Overview**

What is Data Science?

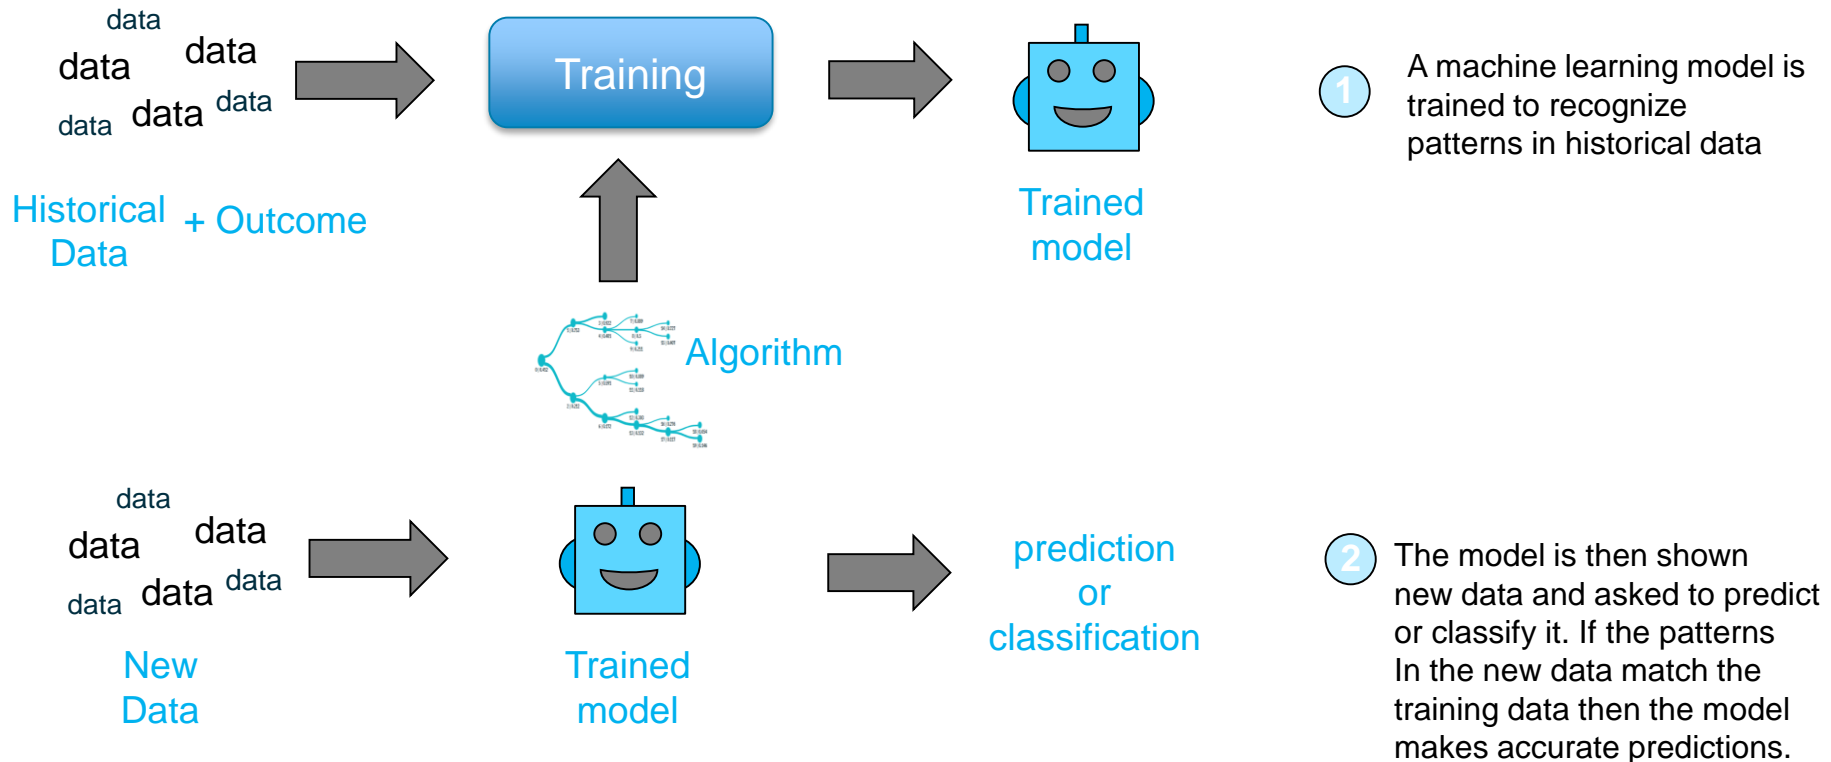


Data Science Projects Require Multiple Skills

Modified from Drew Conway's Venn Diagram

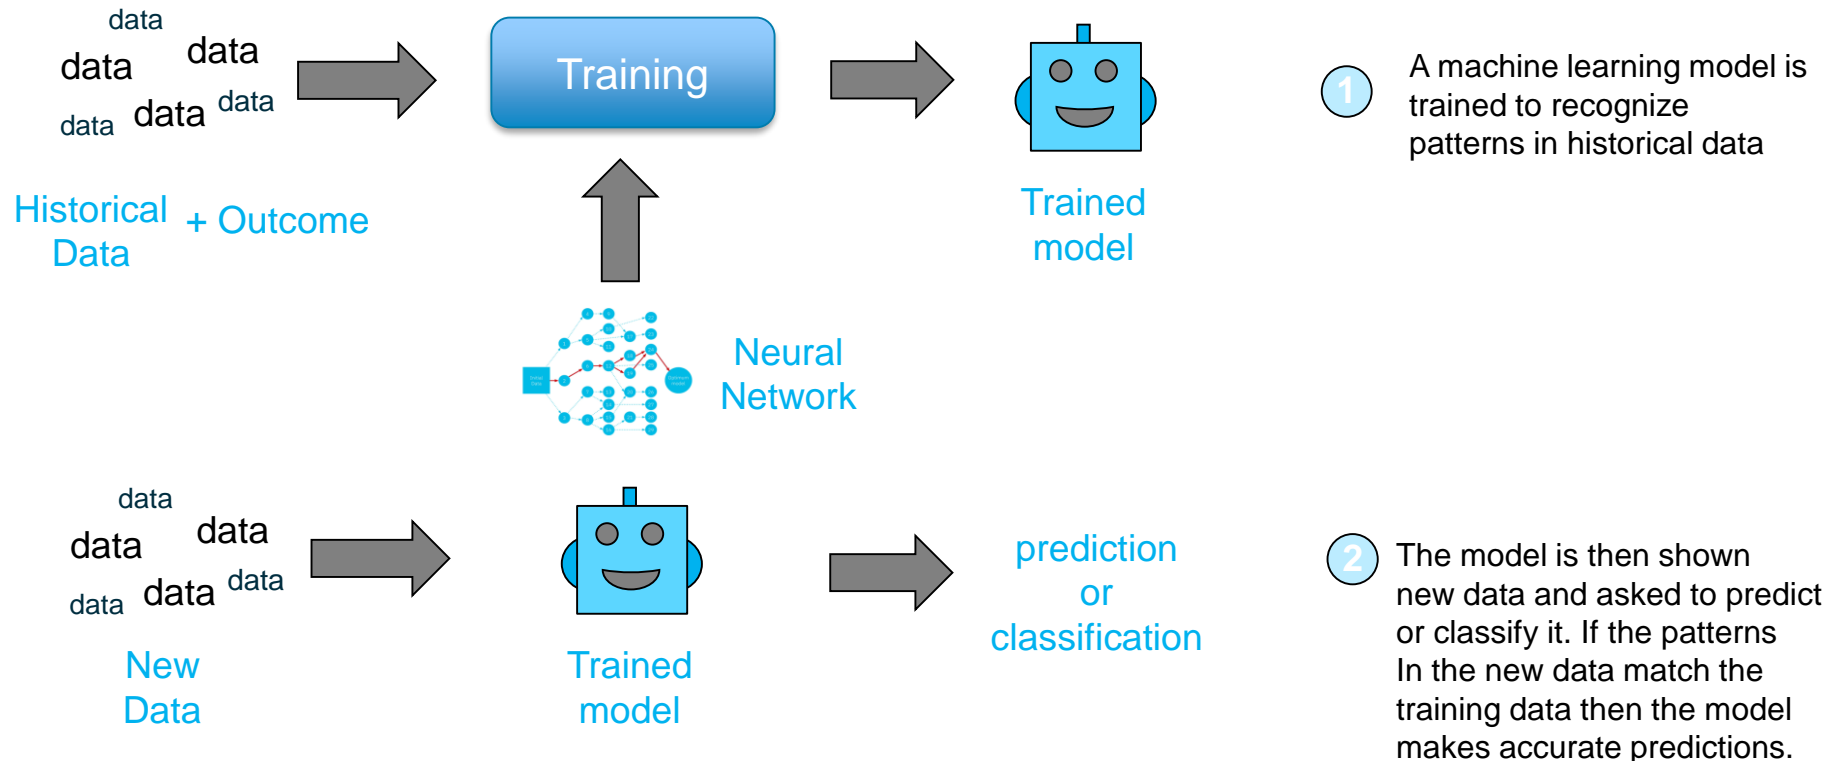
What is Machine Learning?

*“Computers that learn without being **explicitly programmed**”*



What is Deep Learning?

*“Computers that learn without being **explicitly programmed**”*



IBM takes an Enterprise Approach to Data Science

- Freedom of Choice
 - Choose programming languages, open source libraries, IBM value-add capabilities
 - Code/Click
 - Machine Learning/Deep Learning/Decision Optimization.
 - All Data
- Operationalize Machine Learning
 - Manage complete ML lifecycle – Build, Deploy, Manage, Scale, Monitor, Re-train
- Hybrid ML
 - Build where you want, deploy where you want
- Governance
 - Ensure that right people get access to the right data


Outline

- Data Science Overview
- Watson Studio Overview
- Lab Overview



IBM Watson Studio Platform

An integrated platform of tools, services, data, and metadata that help companies or agencies accelerate their shift to be data-driven organizations.



Watson Studio Deployment Options

- Watson Studio on IBM Cloud
 - Managed offering provided by IBM

- Watson Studio Desktop

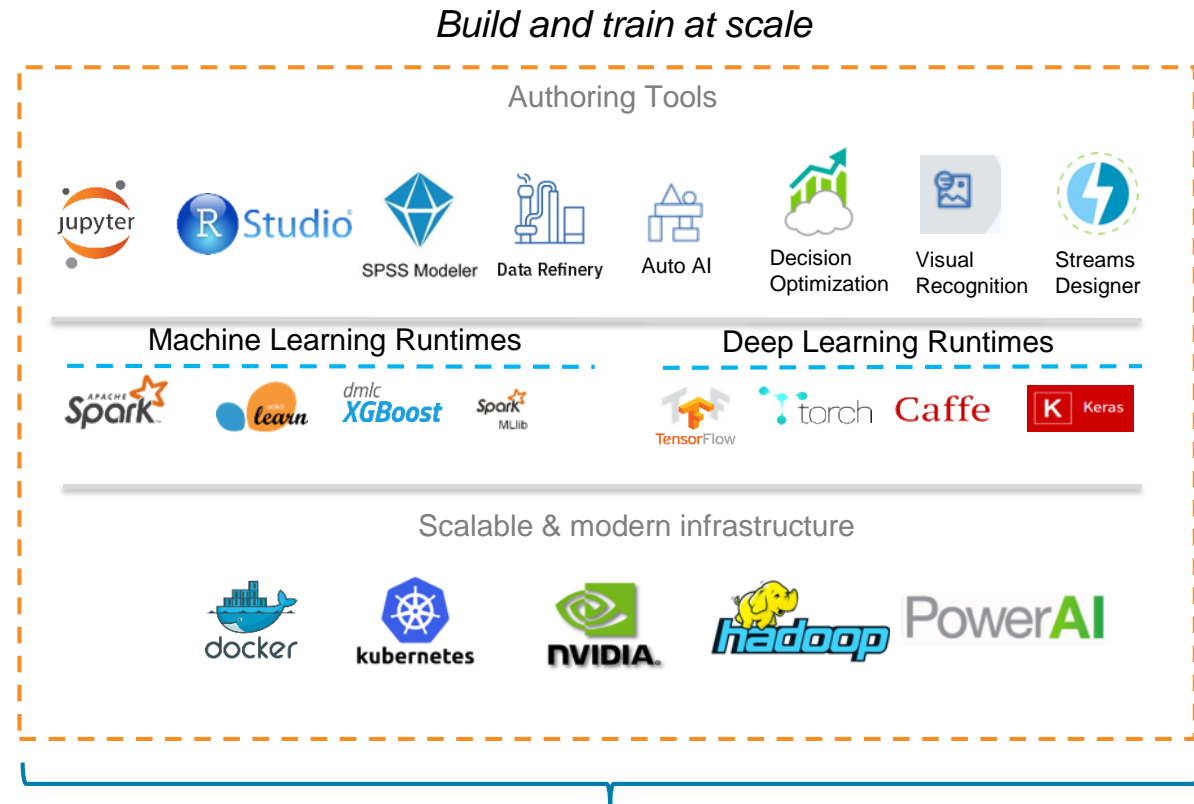
- IBM Cloud Private for Data
 - Watson Studio Local

Watson Studio Tools

- Using best of breed - Open source & IBM tools
- Code (R, Python or Scala) and no-code/visual modeling tools

- Container-based resource management
- Elastic cpu/gpu power
- Run on x86, Power, zLinux
- Integrate with Hadoop/Spark Infrastructure

- Train and deploy where your data lives



IBM
IBM Cloud
Fully Managed


On-prem

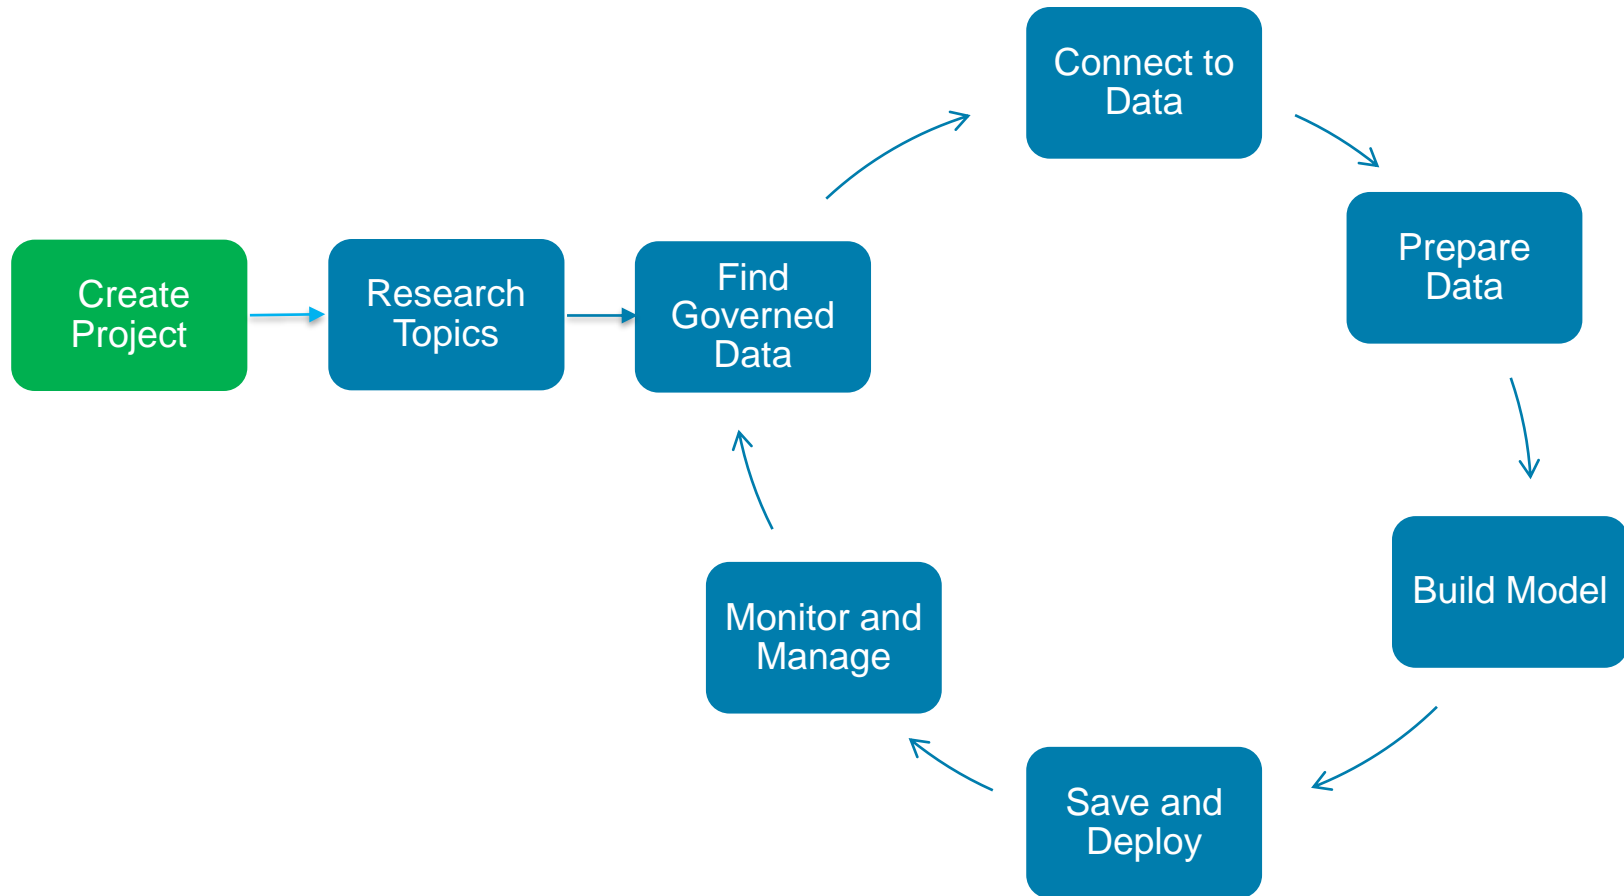
IBM
IBM Cloud

 **amazon**
web services

 **Azure**

Watson Studio supports the Data Science Lifecycle

Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.



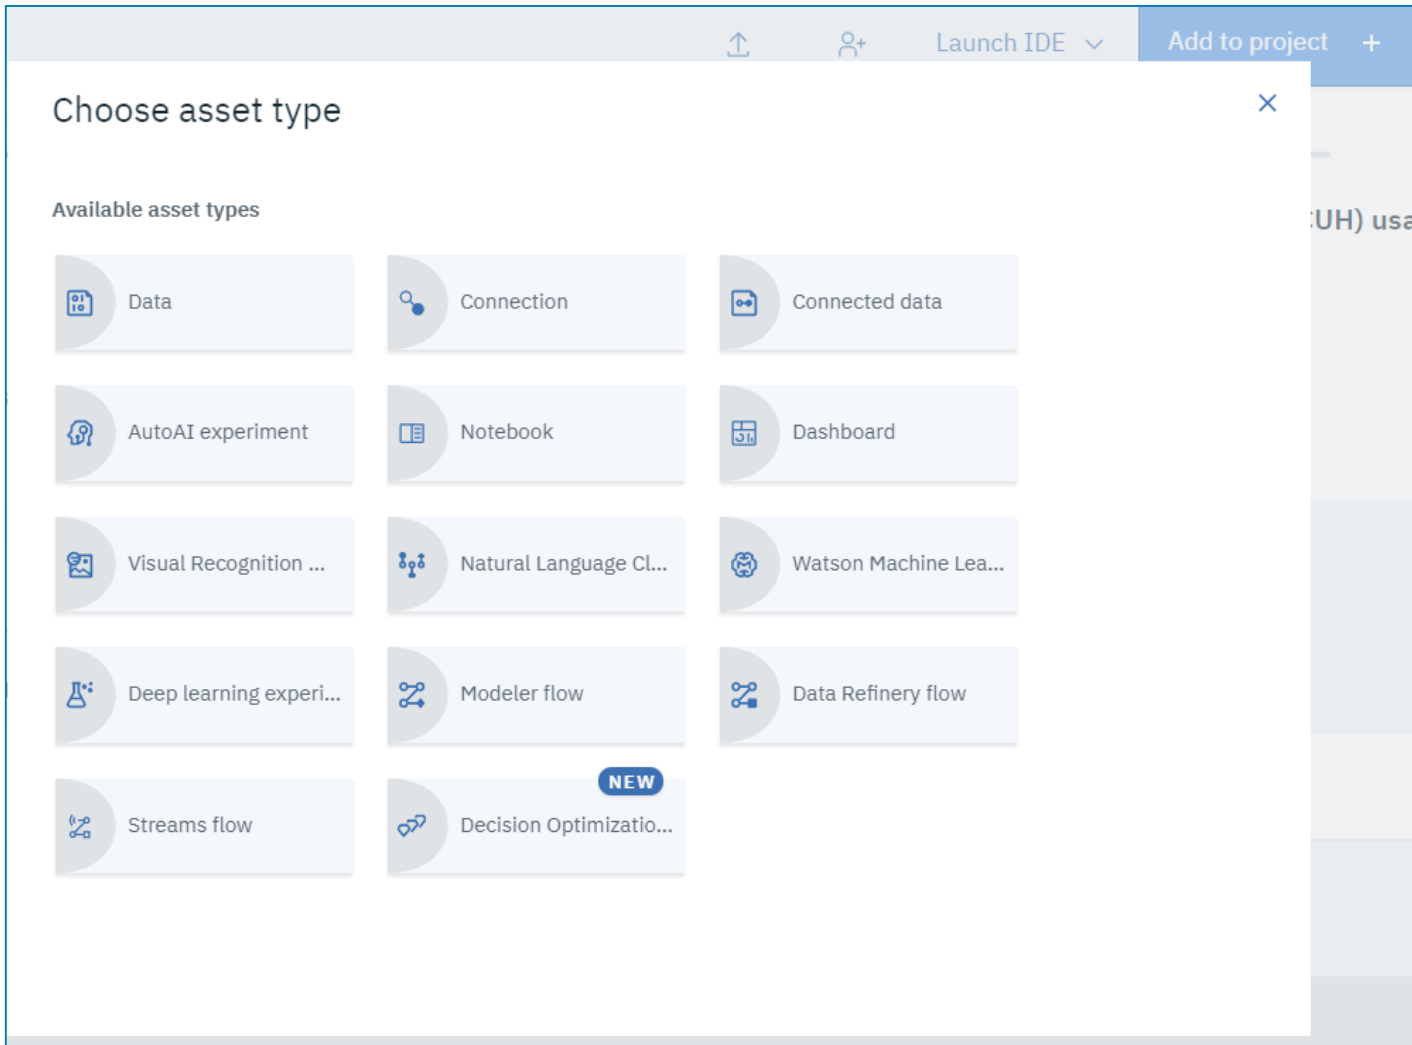
Watson Studio Project Features

Making Data Science a Team Sport

Create
Project















- Organizes resources to achieve a particular data analysis goal
- Support role-based collaboration (Admin, Editor, Viewer)
- Assets from all IDEs can be included in one Watson Studio project: notebooks, data sources, flows, models, etc.
- Export/Import Projects

Add to Project




Choose asset type

Available asset types

 Data	 Connection	 Connected data
 AutoAI experiment	 Notebook	 Dashboard
 Visual Recognition ...	 Natural Language CL...	 Watson Machine Lea...
 Deep learning experi...	 Modeler flow	 Data Refinery flow
 Streams flow	 Decision Optimizatio... NEW	

Watson Studio supports the Data Science Lifecycle

Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.



Watson Studio Gallery

Built-in learning to get started

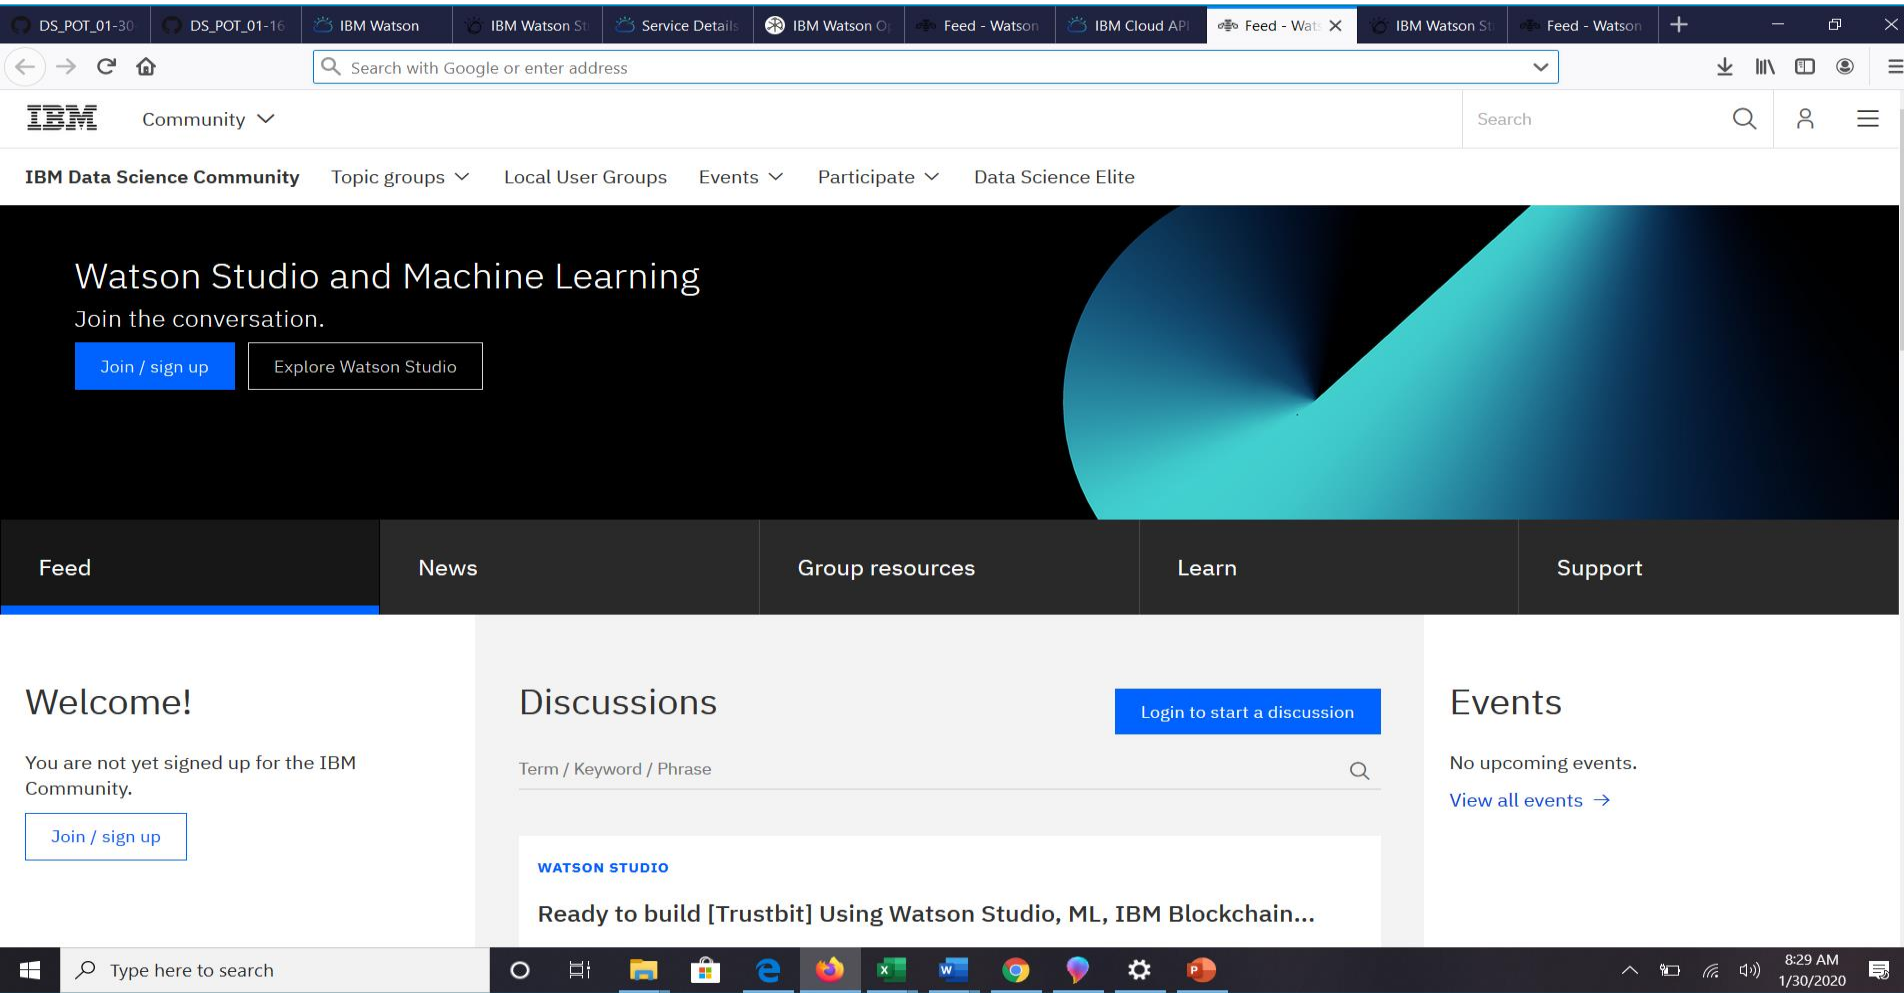
Research
Topics

- The Gallery includes notebooks, and data sets
- Copy notebooks or Data Sets into projects
- Continuously updated in IBM's managed service

Watson Studio Community



[Research Topics](#)

DS_POT_01-30DS_POT_01-16IBM WatsonIBM Watson StService DetailsIBM Watson OFeed - WatsonIBM Cloud APIFeed - WatXIBM Watson StFeed - Watson+

Search with Google or enter address

IBMCommunity ▾

Search



IBM Data Science CommunityTopic groups ▾Local User GroupsEvents ▾Participate ▾Data Science Elite

Watson Studio and Machine Learning

Join the conversation.

[Join / sign up](#)[Explore Watson Studio](#)

Feed

News

Group resources

Learn

Support

Welcome!


You are not yet signed up for the IBM Community.

[Join / sign up](#)

Discussions

[Login to start a discussion](#)

Term / Keyword / Phrase



WATSON STUDIO


Ready to build [Trustbit] Using Watson Studio, ML, IBM Blockchain...


Events

No upcoming events.

[View all events →](#)

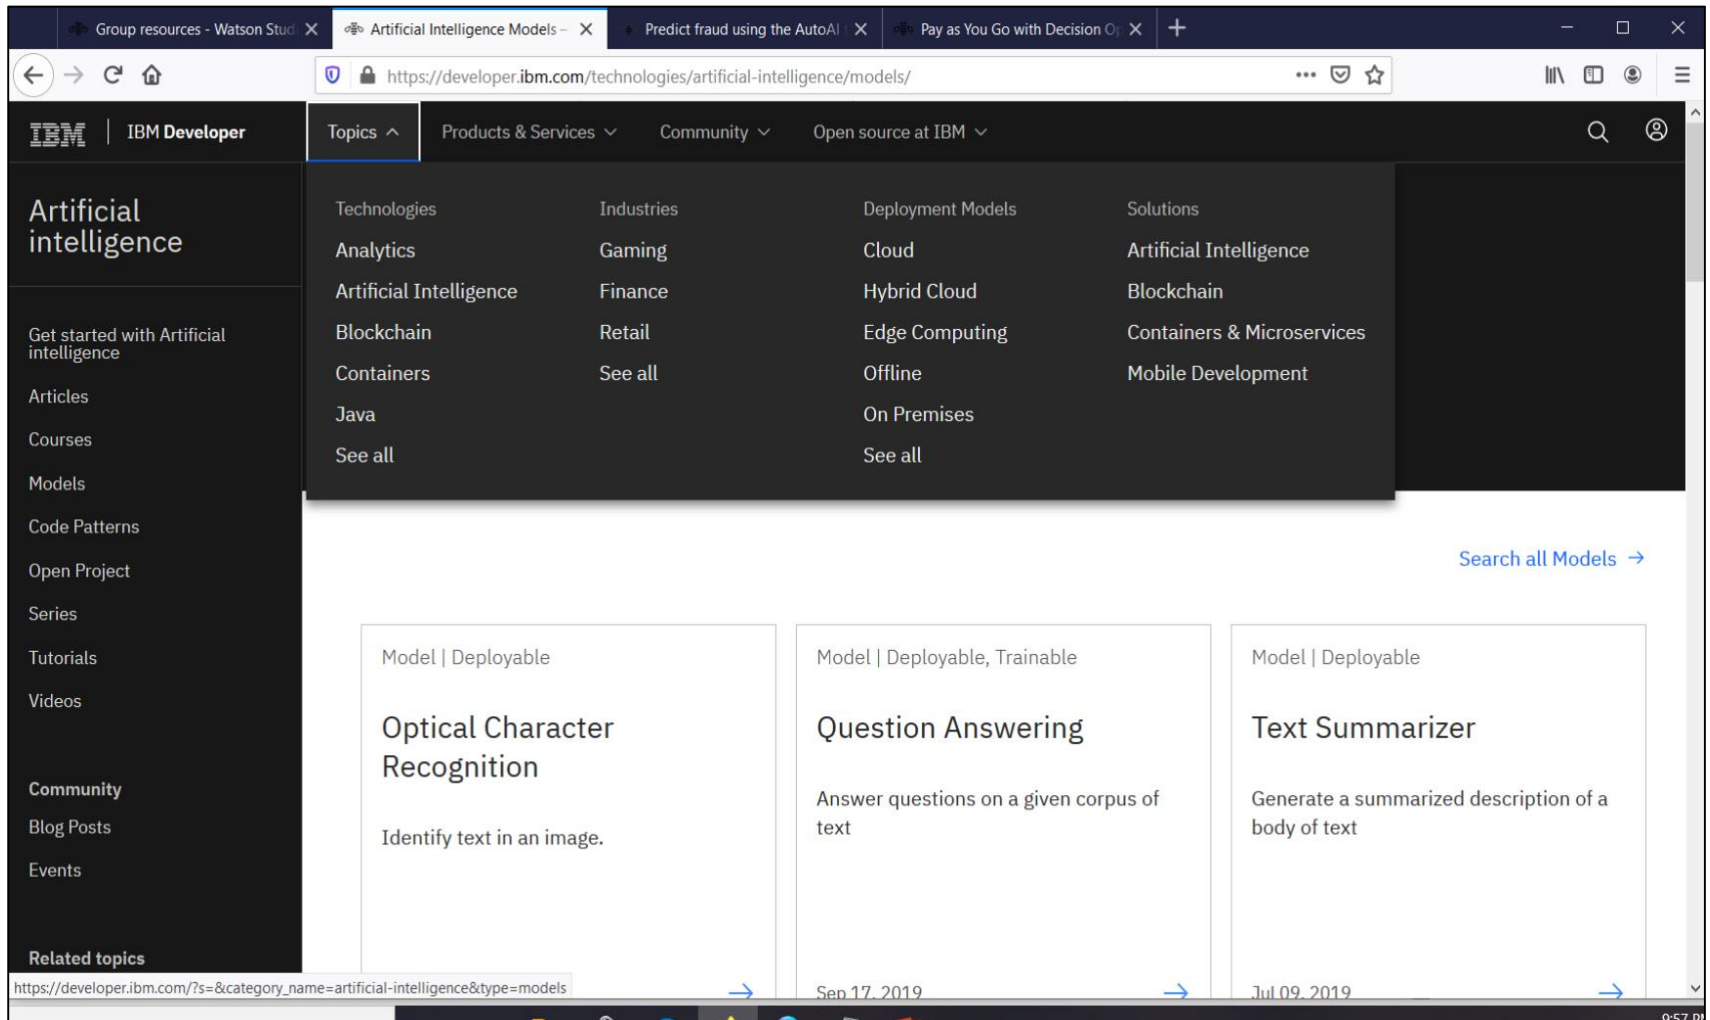
Type here to search



8:29 AM
1/30/2020

developer.ibm.com

Research
Topics



The screenshot displays the IBM Developer website's 'Artificial Intelligence Models' page. The browser's address bar shows the URL <https://developer.ibm.com/technologies/artificial-intelligence/models/>. The page features a dark navigation bar with the IBM logo and 'IBM Developer' text. Below this, a 'Topics' dropdown menu is open, displaying a grid of categories including Technologies, Industries, Deployment Models, and Solutions. The left sidebar contains links for 'Artificial intelligence', 'Get started with Artificial intelligence', 'Articles', 'Courses', 'Models', 'Code Patterns', 'Open Project', 'Series', 'Tutorials', 'Videos', 'Community', 'Blog Posts', 'Events', and 'Related topics'. The main content area shows three model cards: 'Optical Character Recognition' (Model | Deployable), 'Question Answering' (Model | Deployable, Trainable), and 'Text Summarizer' (Model | Deployable). Each card includes a brief description and a date. The bottom of the page shows a search bar and a timestamp of 9:57 PM.

Group resources - Watson Stud X Artificial Intelligence Models X Predict fraud using the AutoAI X Pay as You Go with Decision O X +

← → ↻ 🏠 🔒 <https://developer.ibm.com/technologies/artificial-intelligence/models/> ... 📄 🌐 ☰

IBM | IBM Developer Topics ^ Products & Services ▾ Community ▾ Open source at IBM ▾ 🔍 👤

Artificial intelligence

Get started with Artificial intelligence

Articles

Courses

Models

Code Patterns

Open Project

Series

Tutorials

Videos

Community

Blog Posts

Events

Related topics

Technologies

Analytics

Artificial Intelligence

Blockchain

Containers

Java

See all

Industries

Gaming

Finance

Retail

See all

Deployment Models

Cloud

Hybrid Cloud

Edge Computing

Offline

On Premises

See all

Solutions

Artificial Intelligence

Blockchain

Containers & Microservices

Mobile Development

Search all Models →

Model | Deployable

Optical Character Recognition

Identify text in an image.

Sep 17, 2019

Model | Deployable, Trainable

Question Answering

Answer questions on a given corpus of text

Sep 17, 2019

Model | Deployable

Text Summarizer

Generate a summarized description of a body of text

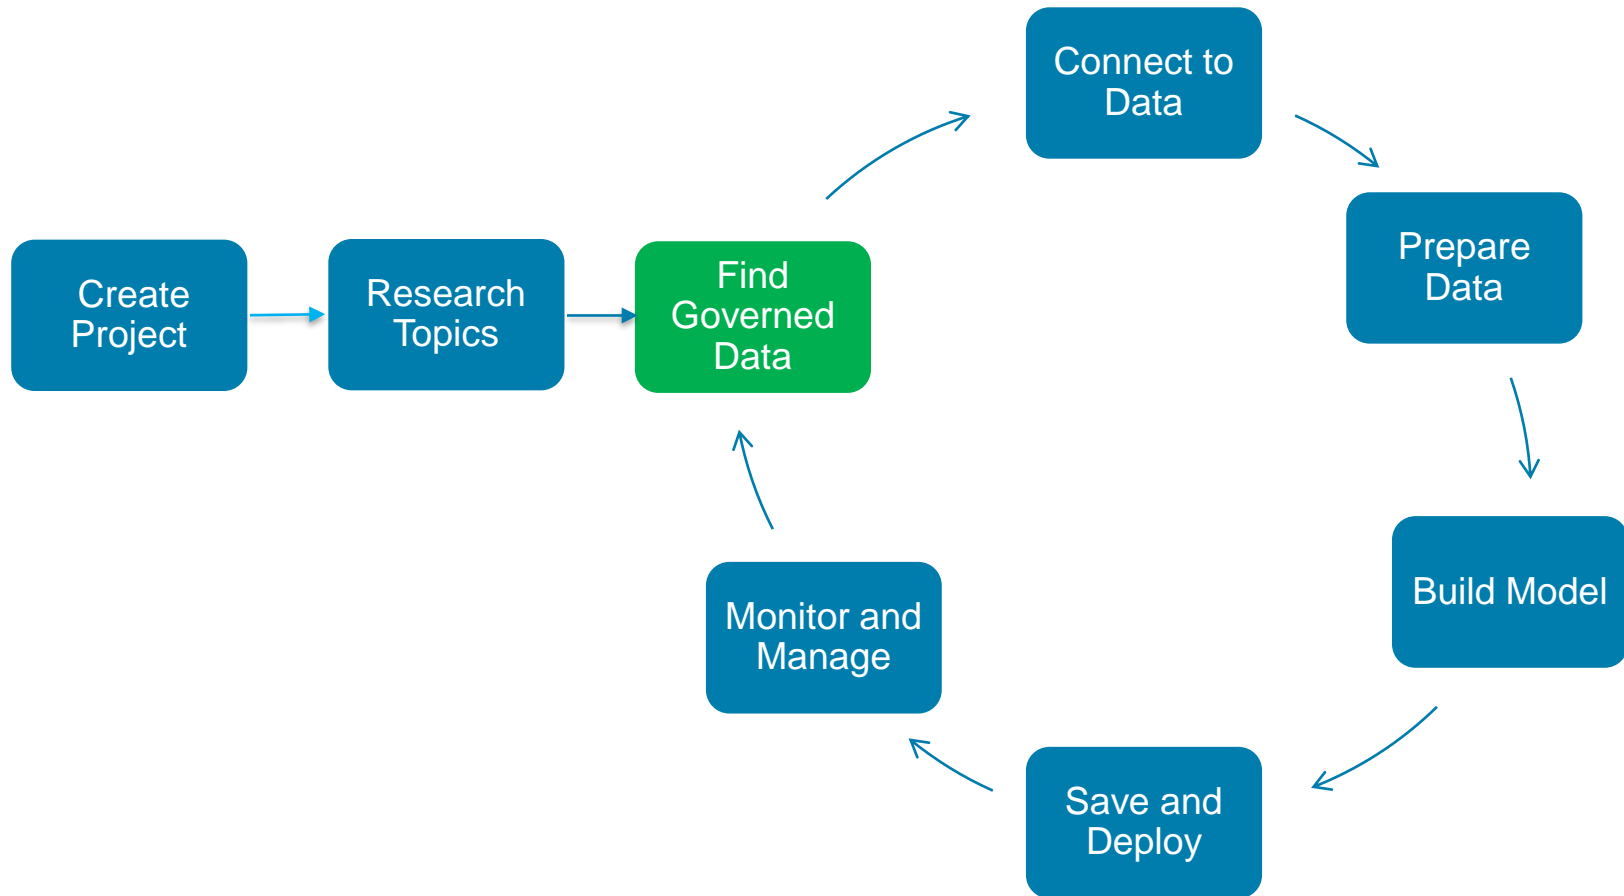
Jul 09, 2019

https://developer.ibm.com/?s=&category_name=artificial-intelligence&type=models →

9:57 PM

Watson Studio supports the Data Science Lifecycle

Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.



Watson Knowledge Catalog Features

Unlock tribal knowledge and unleash knowledge workers

Find
Governed
Data

- **Find** data (structured, unstructured) and AI assets (e.g., ML/DL models, notebooks, Watson Data Kits) in the **Knowledge Catalog** with intelligent search and giving the right access to the right users.
- Discover assets, profiling, classification
- Policy, rule authoring
- Policy, rule enforcement
- Asset Usage Statistics

Watson Knowledge Catalog Features

Find
Governed
Data

Data Asset

female_human_trafficking

Description

There is no description available for this asset.

Added: Jan 31, 2019 10:02 AM
Format: application/octet-stream
Size: 347 KB

Tags

trafficking | female human trafficking

Reviews

☆☆☆☆☆ 0 reviews

Connection

Source: [Watson Studio Labs_DataCatalog](#)
Source type: Cloud Object Storage

Classification

Personally Identifiable Information
Personally identifiable information (PII) is defined as any data that could potentially identify a specific individual. Any information that can be used to distinguish one person from another can be considered PII.

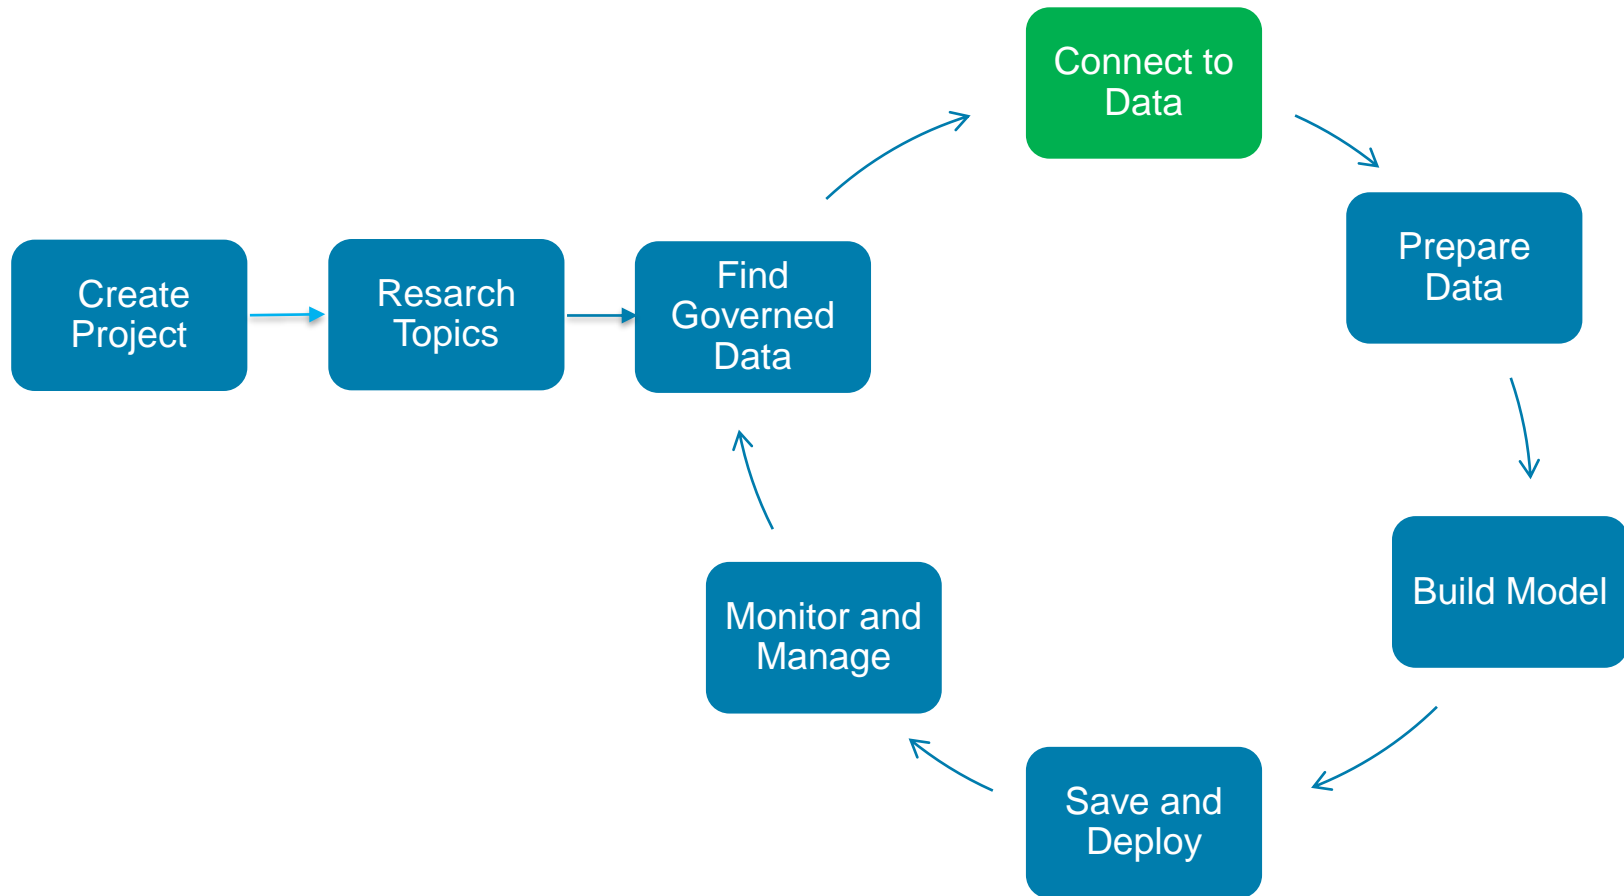
Schema: 26 Columns | 1085 Rows | 2 Columns anonymized ⓘ

Preview: 1000 rows | Last refresh: 22 seconds ago | [Refresh](#)

DATE	BIRTH_COUNT...	BIRTH_COUNTRY_CODE	OCCUPATION	ADDRESS	SSN	PASSPORT_NUMBER
	Type: String	Type: String	Type: String	Type: String	Type: String	Type: String
th	Country Name	Country Code	Text	Text	US Social ...	Passport Number
15	Ghana	GH	Engineer, land	824 Kristin Grv, A	afe55d1d355c3:	1c9da91e1e20863dd850
19	Ghana	GH	Editor, commissi	1148 Wang Fall S	77a0daa42ec7d	12d38855ed107e7cc5dd
16	Ghana	GH	Merchant navy of	9486 Pratt Wall,	669061087d6d1	c43ed0283a3def7031d8:
17	Ghana	GH	Paramedic	0890 Johnson Tr	997b59e501b2e	179abee5ba608418154d
18	Ghana	GH	Surveyor, buildin	2315 Brittany Cr	70329b83b40cb	84524ccc3c5c6590600e:
24	Ghana	GH	Waste managem	88811 Donald Pa	d2f2236f52407:	a730ae13f5ed96f71e904
23	Ghana	GH	Doctor, general p	9150 Donald Rpe	d2c2d41163d8f:	ced1617be1d70e44421c
02	Ghana	GH	Forest/woodland	1355 Lopez Villa	62007942c2b0c	8c8debda401b6b6d954b
12	Ghana	GH	Land/geomatics :	86792 Amy Vlgs,	08f8dd9f9ba89t	a43f1d6c9cacdfa82a1a1
10	Ghana	GH	Oncologist	108 Erin Via, Nev	f8b871f6e058e2	f289be62078ebbe457c6:
07	Ghana	GH	Veterinary surge	79572 Schmidt E	f2006c1d30df33	624a9605774a0cfd98aa

Watson Studio supports the Data Science Lifecycle

Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.



Watson Studio Connection Features

Connect to
Data



















- Upload files
- Connectors to Structured and Unstructured, On-prem and Cloud data sources.
- Wizard based connection definition and code generation

Connection Options























Connect to
Data

New connection

IBM services

 BigInsights HDFS	 Cloud Object Storage	 Cloud Object Storage (infrastructure)	 Cloudant
 Cognos Analytics	 Compose for MySQL	 Compose for PostgreSQL	 Db2
 Db2 Big SQL	 Db2 for i	 Db2 for z/OS	 Db2 Hosted
 Db2 on Cloud	 Db2 Warehouse	 Informix	 Object Storage OpenStack Swift (infrastructure)
 PureData for Analytics	 Watson Analytics		

Third-party services

 Amazon Redshift	 Amazon S3	 Apache Hive	 Cloudera Impala
 Dropbox	 FTP	 Google BigQuery	 Google Cloud Storage
 Hortonworks HDFS	 Looker	 Microsoft Azure Data Lake Store	 Microsoft Azure SQL Database
 Microsoft SQL Server	 MySQL	 Oracle	 Pivotal Greenplum
 PostgreSQL	 Salesforce.com	 Sybase	 Sybase IQ
 Tableau	 Teradata		

Notebook Screenshot

[Connect to Data](#)

The screenshot shows the IBM Analytics Notebook interface. The top menu bar includes File, Edit, View, Insert, Cell, Kernel, Widgets, and Help. The right side of the top bar shows 'Trusted', 'Python3.6', and tabs for 'Local', 'Remote', and 'Other'. Below the menu bar is a toolbar with icons for file operations, running, and markdown. The main area contains a code cell with a list of database schema fields. Below the code cell is a section titled 'Read in CLAIMS Table'. The bottom code cell shows a Python script for connecting to a remote data source and loading data into a Spark DataFrame. The right sidebar displays a list of data sources: PROCEDURES, PATIENTS, and CLAIMS, each with an 'Insert to code' button. A dropdown menu is open for the 'CLAIMS' source, showing options to 'Insert Pandas DataFrame' or 'Insert Spark DataFrame in Python'.

```
-- PROCEDURE_PERCT_RANK: integer (nullable = true)
-- PROCEDURE_RISK_GROUP: string (nullable = true)
-- QUANTITY_INDEX: integer (nullable = true)
-- SERVICE_TYPE: string (nullable = true)
-- SUBMIT_CHG: integer (nullable = true)
-- SUBMITTED_CHG_INDEX: integer (nullable = true)
-- TOTAL_CHARGES_INDEX: integer (nullable = true)
-- TOTAL_CHARGES_PER_PROCEDURE: integer (nullable = true)
-- USER_DEFINED_FLAG_0: string (nullable = true)
-- SUBMITTED_CHARGE_AMOUNT: integer (nullable = true)
-- CLAIM_NUMBER: string (nullable = true)
-- IS_FRAUD: string (nullable = true)
```

Read in CLAIMS Table

In [5]:

```
import dsx_core_utils, requests, os, io
from pyspark.sql import SparkSession
# Add asset from remote connection
df7 = None
dataSet = dsx_core_utils.get_remote_data_set_info('CLAIMS')
dataSource = dsx_core_utils.get_data_source_info(dataSet['datasource'])
sparkSession = SparkSession(sc).builder.getOrCreate()
# Load JDBC data to Spark dataframe
dbTableOrQuery = ''' + (dataSet['schema'] + '`.`' if len(dataSet['schema'].strip()) != 0 else '') + dataSet['table'] + '''
if (dataSet['query']):
```

PROCEDURES
Insert to code

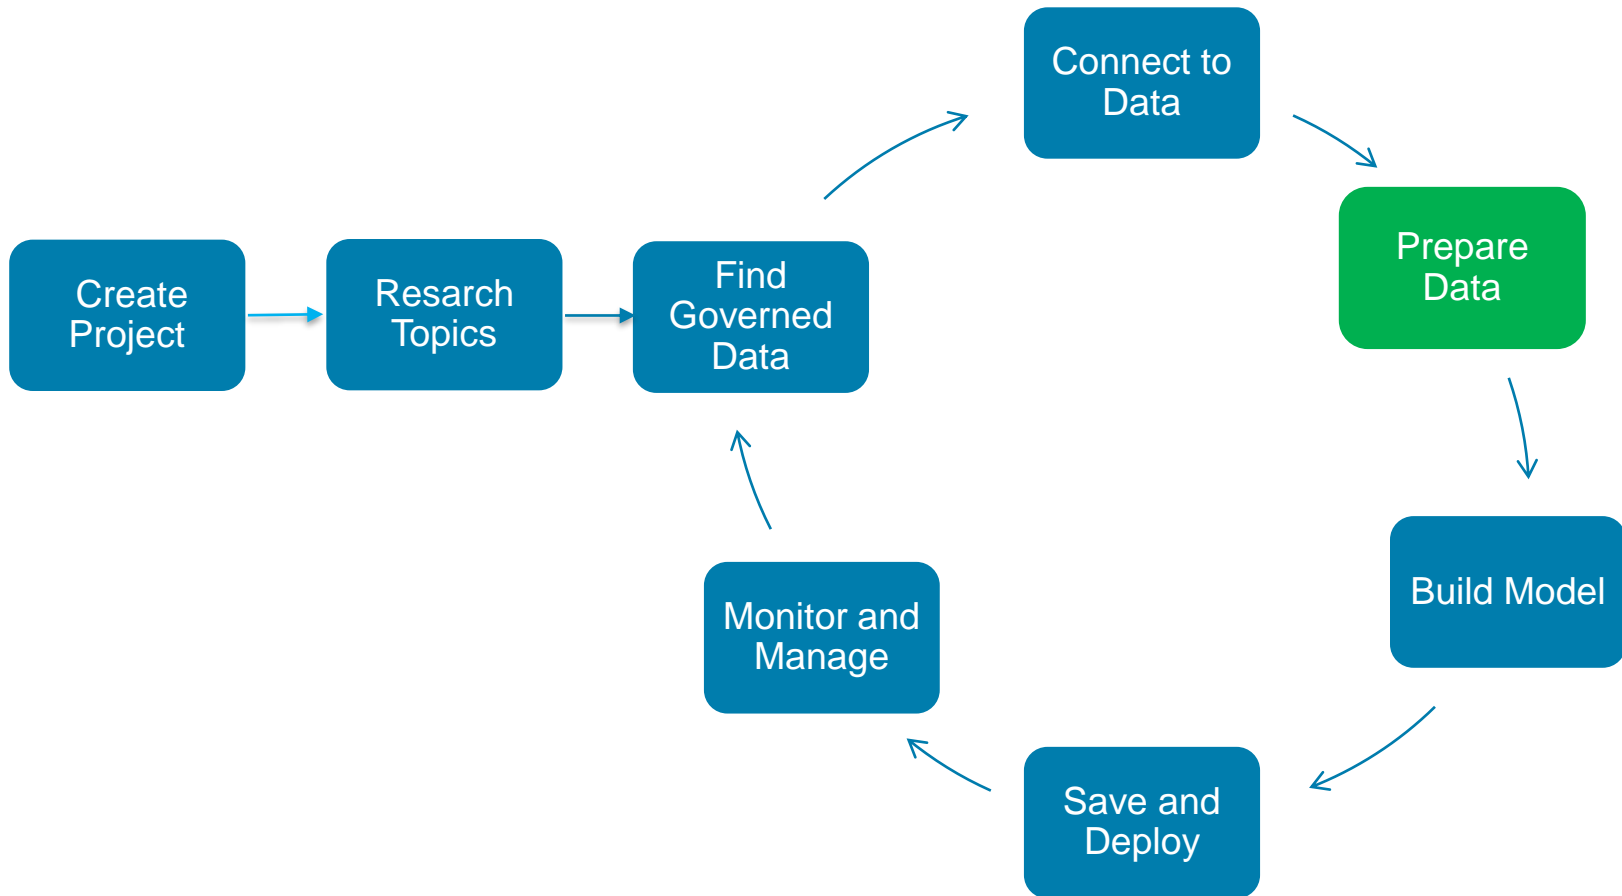
PATIENTS
Insert to code

CLAIMS
Insert to code

- Insert Pandas DataFrame
- Insert Spark DataFrame in Python

Watson Studio supports the Data Science Lifecycle

Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.



Watson Studio Data Refinery Features

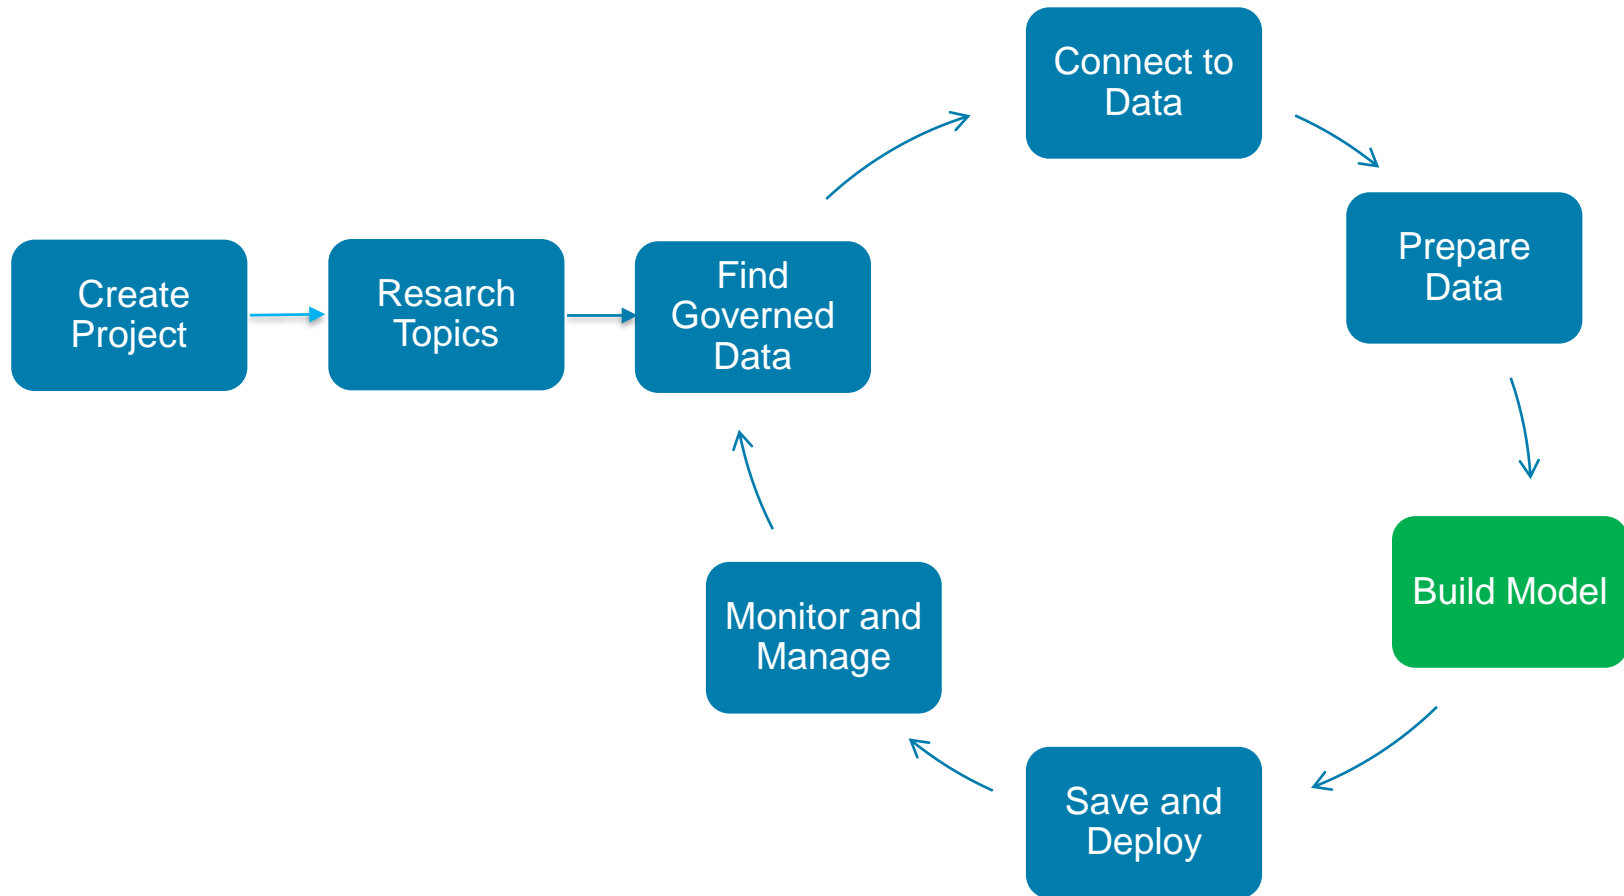
Making Data fit for use

Prepare
Data

- Data Refinery tool to profile, visualize, and shape data.
- Create data preparation pipelines via point and click capability on subset of data
 - ✓ Cleanse the data: fixing or removing data that is incorrect, incomplete, improperly formatted, or duplicated
 - ✓ Shape the data: customize data by filtering, sorting, combining, or removing columns, and performing operations
- Run the pipeline on all the data
 - Manually (on demand)
 - Automated (scheduled)

Watson Studio supports the Data Science Lifecycle

Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.



Watson Studio Model Building Features

The best of open source and IBM Watson tools to create start-of-the-art data products

Build Model

Open Source Tools

- Jupyter Notebooks**
- RStudio and Shiny
- Libraries- scikit-learn**, XGBoost**, Spark**, TensorFlow, Caffe, Keras, PyTorch

IBM Tools

- AutoAI **
- SPSS Modeler**
- Neural Network Modeler
- Experiment Builder
- Natural Language Classifier Model
- Visual Recognition Model


Train at scale on **GPUs** and **distributed** compute

** in hands-on labs

Jupyter Notebook

[Build Model](#)

My Projects / Watson Studio Labs / Machine Learning with SparkML

File Edit View Insert Cell Kernel Help

Not Trusted | Python 3.6 with Spark        Run   Format Markdown 

Read Data Asset - female_human_trafficking - See Lab Instructions

```
In [ ]: # Insert SparkSession DataFrame code in this cell after the comments.  
# make CERTAIN to rename the default dataframe name (df_data_1 or df_data_2 or df_data_3, etc) to trafficking_df  
# Put cursor on the next line to Insert to code.
```

Read Data Asset - Occupations - See Lab Instructions

The occupations listed in the female human trafficking file are too numerous to use as input to a machine learning model. We will categorize these occupations into 15 categories by joining with two other files. The Occupation.csv file contains a mapping of the occupations in the female human trafficking table to a category code. The Categories.csv file contains each code followed by the category name. This information needs to be joined to the female human trafficking table.

Follow the same procedure as above to insert a SparkDataFrame for Occupations

```
In [ ]: # Insert SparkSession DataFrame code in this cell after the comments  
# make CERTAIN to rename the default dataframe name (df_data_1 or df_data_2 or df_data_3,etc) to occupations  
#Put cursor on the next line to Insert to code
```

Read Data Asset - Categories - See Lab Instructions

Follow the same procedure as above to insert a SparkDataFrame for Categories

```
In [ ]: # Insert SparkSession DataFrame code in this cell after the comments  
# make CERTAIN to rename the default dataframe name (df_data_1 or df_data_2 or df_data_3,etc) to categories  
#Put cursor on the next line to Insert to code
```


SPSS Modeler

[Build Model](#)

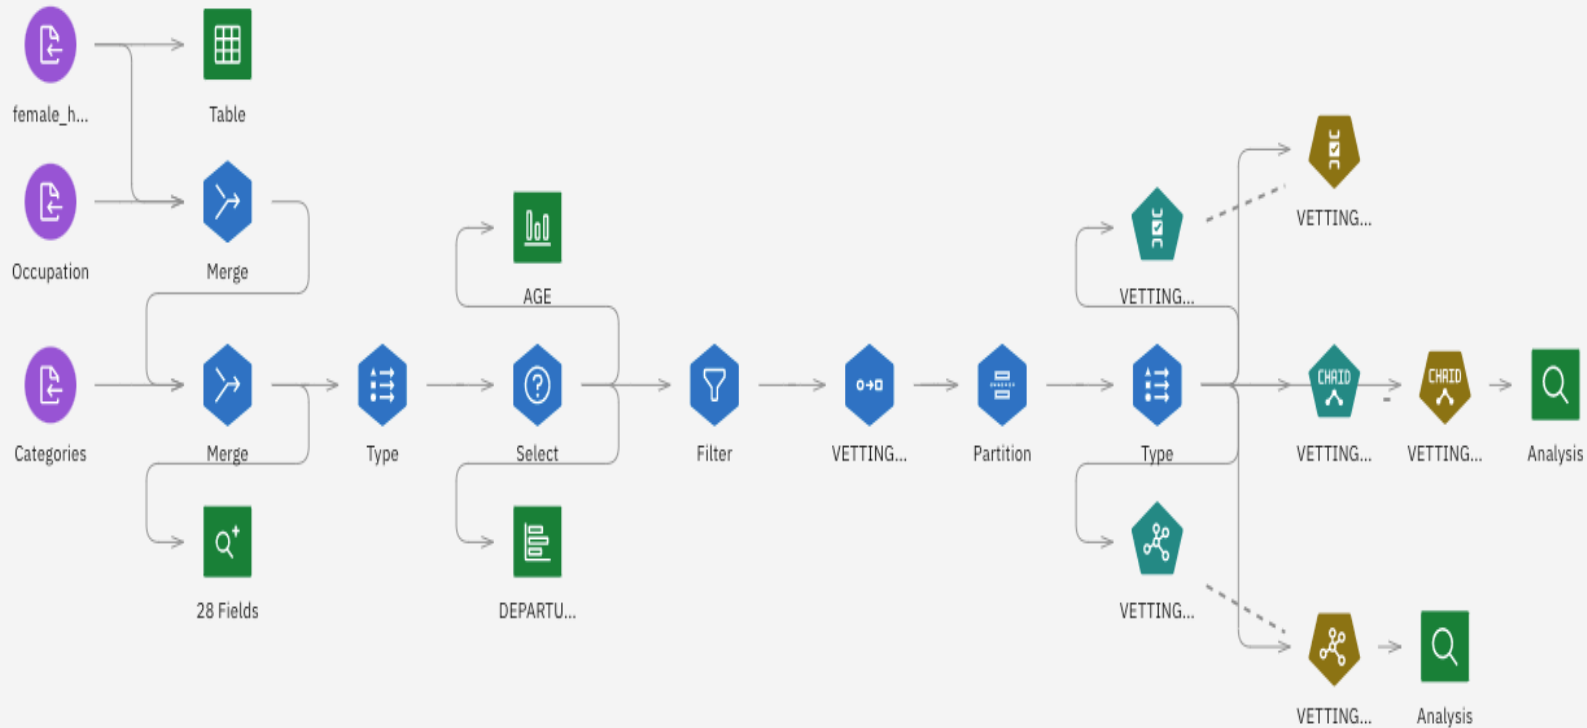
My Projects / Watson Studio Labs / Titanic SPSS



Search Palette



- Import
- Record Operations
- Field Operations
- Graphs
- Modeling
- Outputs
- Export



Neural Network Modeler

[Build Model](#)

An intuitive drag-and-drop, no-code interface for designing neural network structures using the most popular deep learning frameworks. Quickly capture your network design then single click export for experimental optimization.


Supported Frameworks TensorFlow™  PYTORCH

Caffe

 Keras

Drag-and-drop
network layers

Real-time validation of network
flow

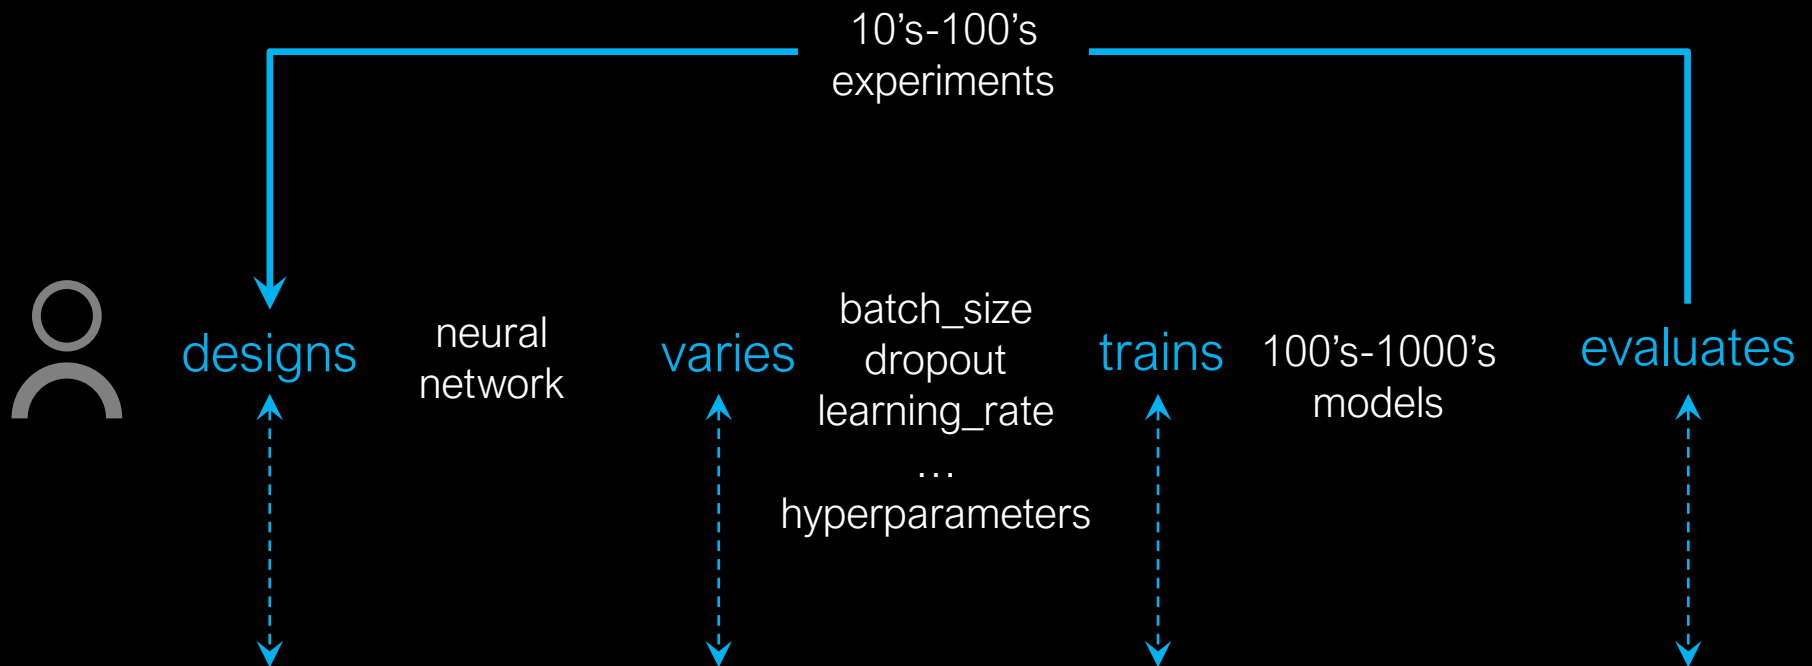


- Define layer configuration
- Choose optimizer params

- Generate CPU or GPU compatible code

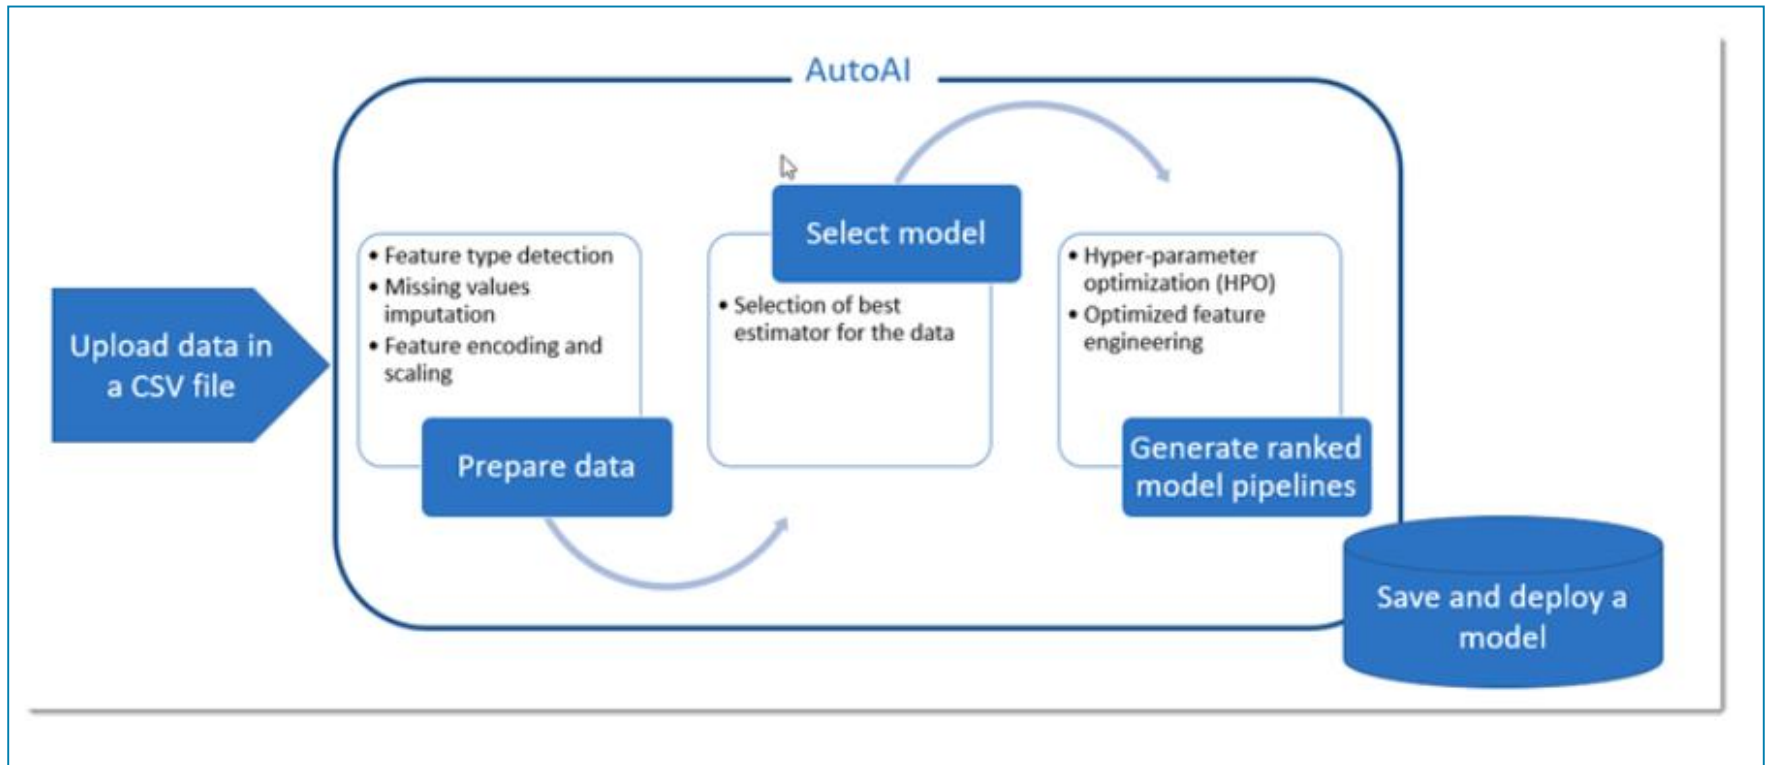
- Save as popular framework code
- Export as a python notebook
- Execute as batch experiment

Experiment Builder

[Build Model](#)


Experiment Builder
supports the end-to-end workflow

AutoAI

Build Model

AutoAI

Build Model



Pipeline leaderboard

[Compare pipelines](#)


Ranking based on:

Accuracy

	RANK	ACCURACY	PIPELINE INFORMATION		
>	1	0.897	P3 - XGB classifier estimator Transformers (8): Preprocessing > Standard scaler > Univariate feature selection > Sine > Univariate feature selection > Tangent > ...	View details	Save as model
>	2	0.884	P1 - XGB classifier estimator Transformers (2): Preprocessing > XGB classifier estimator	View details	Save as model
>	3	0.884	P2 - XGB classifier estimator Transformers (2): Preprocessing > XGB classifier estimator	View details	Save as model

Watson Studio supports the Data Science Lifecycle

Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.



Watson Studio Save and Deploy Models

Save and Deploy Models with Watson Machine Learning

Save and
Deploy



dataplatform.cloud.ibm.com/ml/models/8e525b3c-cd00-48a9-9701-e9208dccee5d/new-deployment?projectid=b9150de9-e8d0-4119-af59-d1a0c8fbf7ee&mlInstanceGuid=be47af33-01c...

IBM Watson Projects Tools Catalog Community Services Manage Support Docs 1449375 - IBM Corpor... BB

Create Deployment

Define deployment details

Name

Model

Description

Deployment description

300

Deployment type

- ☒ Web service
☐ Batch prediction
☐ Realtime streaming prediction


Cancel

Save



Watson Studio Save and Deploy Trained Models

Save and Deploy Models with Watson Machine Learning



Watson Studio Save and Deploy Features

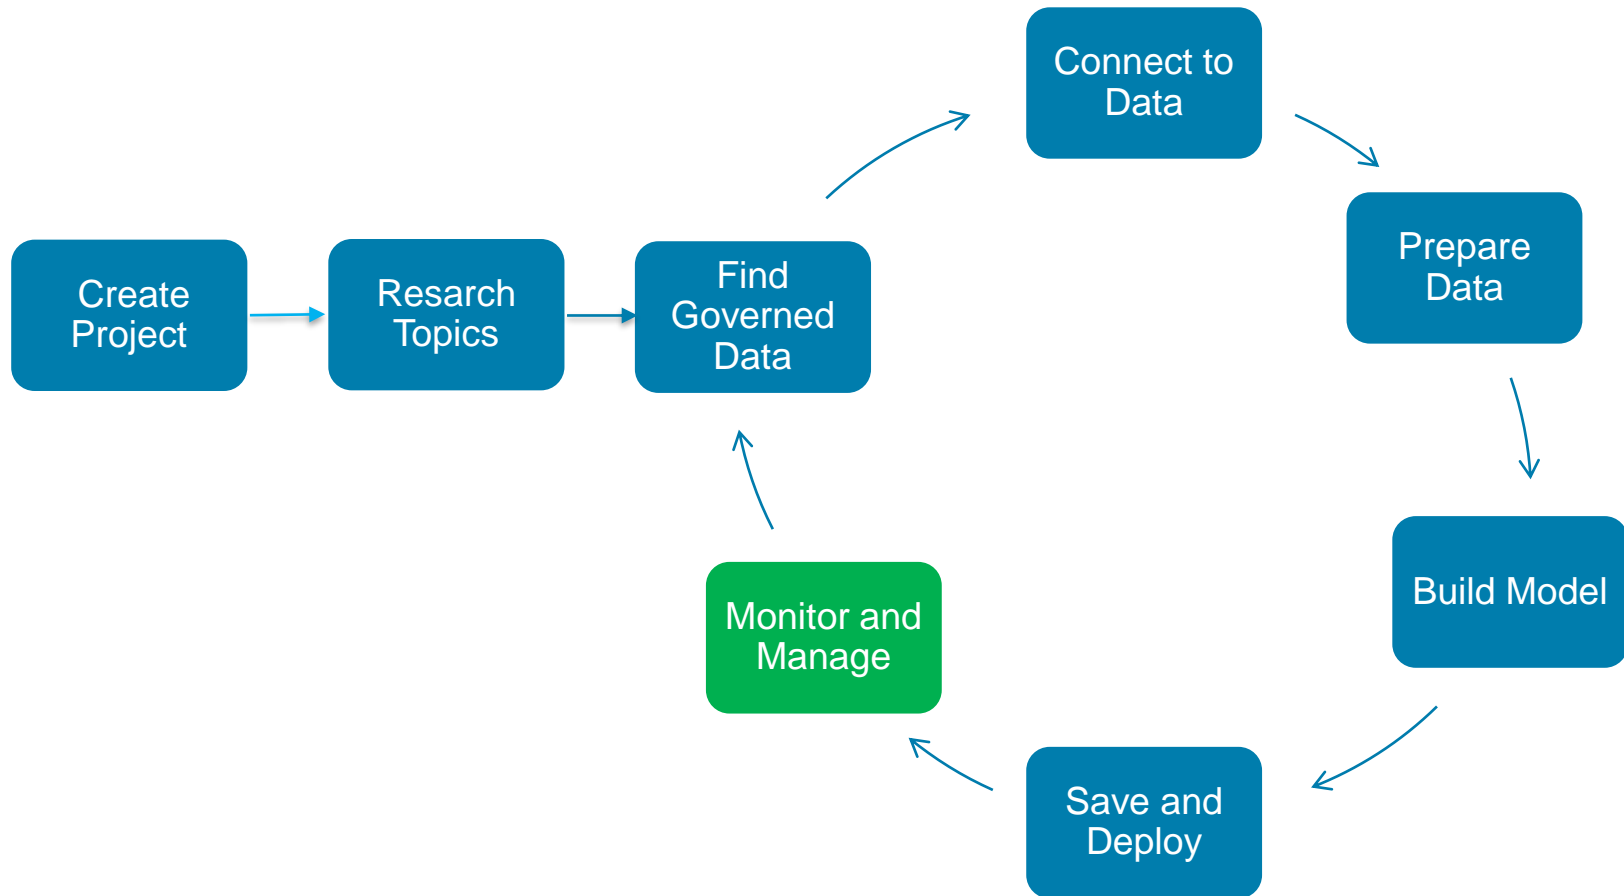
Save and Deploy Models with Watson Machine Learning

Save and
Deploy

- Watson Machine Learning API to save/load models to/from repository
- Watson Machine Learning API to deploy saved models easily and have them scale automatically.
- Watson Machine Learning API to invoke deployed models

Watson Studio supports the Data Science Lifecycle

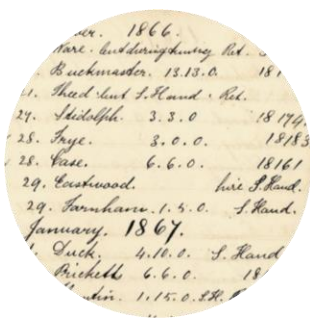
Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.



Our vision for Trusted AI

Pillars of trust, woven into the lifecycle of an AI application


Monitor and
Manage



**Is it
accurate?**



**Is it
fair?**



**Is it easy to
understand
?**



**Did anyone
tamper
with it?**

Watson OpenScale

Monitor and
Manage

Trust and Transparency

- Intelligently delivers bias mitigation help
- Provides traceability & auditability of AI predictions made in production applications
- Tracks AI accuracy in applications
- Explains an outcome in business terms
- Drift Detection

Automation


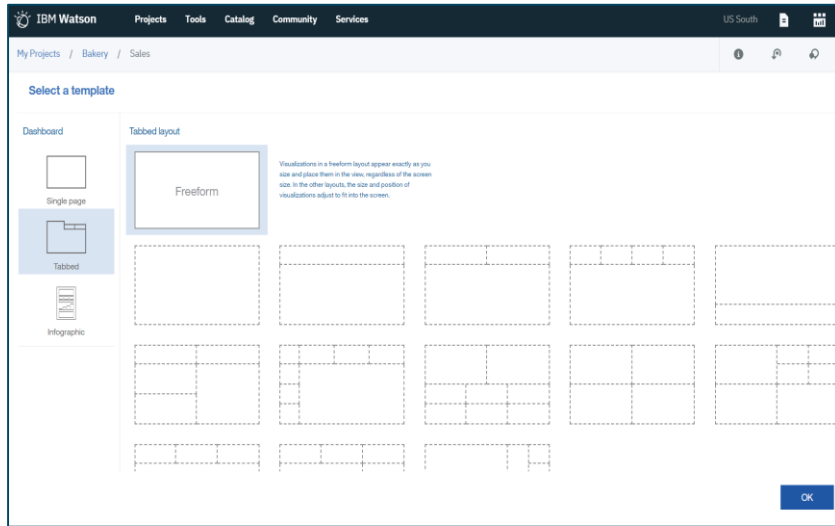
- Automatically detects and mitigates bias in model output, without affecting currently deployed model or outcomes

Open By Design

- Monitor models deployed on third party model server engines
- Deploy behind enterprise firewall or on IaaS provider

Watson Studio Dynamic Dashboards

Making insights available to all



My Projects / Bakery Sales

Select a template

Dashboard

Tabbed layout

Freeform

Visualizations in a freeform layout appear exactly as you size and place them in the view, regardless of the screen size. In the other layouts, the size and position of visualizations adjust to fit into the screen.

OK

My Projects / Bakery Sales

Data assets

0 assets selected.

NAME	TYPE	SERVICE	CREATED BY	LAST MODIFIED	ACTIONS
UNdata_agri_value_add.csv	Data Asset	Project	Alex Jones	7 Mar 2018, 9:37:13 am	
EuropeanCountryStats.csv	Data Asset	Project	Alex Jones	7 Mar 2018, 9:37:12 am	
Bakery Sales and Weather data.csv	Data Asset	Project	Alex Jones	8 Feb 2018, 3:07:05 pm	

Notebooks

NAME SHARED SCHEDULED STATUS LANGUAGE LAST EDITOR LAST MODIFIED ACTIONS

Sales Predictions Alex Jones 7 Mar 2018

Streams flows

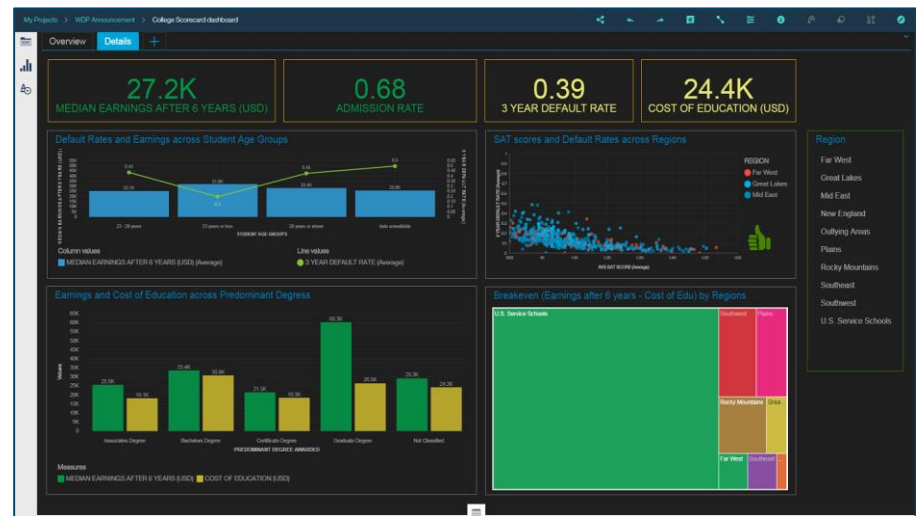
Dashboard

0 assets selected.

NAME	SHARED	LAST EDITOR	LAST MODIFIED	ACTIONS
Bakery Dashboard		Alex Jones	9 Feb 2018, 4:58:46 pm	

Models

Share Remove



Watson Studio Takeaways

Integrated Collaboration Environment

- Data Scientists, Subject Matter experts, Business Analysts & Developers all in one environment to accelerate innovation, collaboration and productivity
- Built-in learning to get started or go the distance with advanced tutorials

Choice of Tools for the full AI lifecycle

- Best in-breed open source and IBM tools that support the end-to-end AI lifecycle
- Choice of code or no-code tools to build and train your own ML/DL models or easily train and customize pre-trained Watson APIs

Support for all levels of expertise

- Use Watson smarts and recommendations for the best algorithms to use given your data, OR
- Use the rich capabilities and controls to fine tune your models

Multiple Deployment Options

- Watson Studio on IBM Cloud – Managed offering
- Watson Studio Local – Private Cloud, Public Cloud-(IBM, Azure, AWS)
- Watson Studio Desktop

Model lifecycle & management

- Deploy models into production then monitor them to evaluate performance.
- Capture new data for continuous learning and retrain models so they continually adapt to changing conditions.



Integrated with Knowledge Catalog

- Intelligent discovery of data and AI assets that enables reuse & improves productivity
- Seamlessly integrated for productive use with Machine Learning and Data science
- Powerful governance tools to control and protect access to data

Outline

- **Data Science Overview**
- **Watson Studio Overview**
- **Lab Overview** 

Breakout Rooms



Share

Leave Breakout Room

Lab Use Case: Female Human Trafficking


Input


- Generated fake travel records based on incoming custom forms.
- Subset of records were vetted as “high”, “medium”, or “low” risk for Female Human Trafficking by an analyst.

Goal is to train a model on the vetted data to be able to score the unvetted travel records into high, medium, or low categories.

Lab Data

Field	Description
UUID	Hash-based unique identifier
VETTING_LEVEL	Analyst vetting status : 100- PENDING, 10 – HIGH, 20 – MED, 10 - LOW
NAME	Person name
GENDER	Person Gender
AGE (SPSS Modeler)	Person age at time of travel
BIRTH_DATE (Notebook)	Person birth date
BIRTH_COUNTRY	Person full birth country
BIRTH_COUNTRY_CODE	Person ISO 2 country
OCCUPATION CATEGORY	Person occupation as declared on form
ADDRESS	Person US address
SSN	Person Social Security Number
PASSPORT_NUMBER	Person Passport Number
PASSPORT_COUNTRY	Person Passport Issuing Country
PASSPORT_COUNTRY_CODE	Person Passport Issuing Country ISO 2 Code
COUNTRYIES_VISITED	The countries visited as declared on form
COUNTRIES_VISITED_COUNT	The number of countries visited as declared on form
ARRIVAL_AIRPORT_COUNTRY_CODE	ARRIVAL Airport country code ISO2
AIRPORT_ARRIVAL_IATA	ARRIVAL Airport 3 character code
AIRPORT_ARRIVAL_MUNICIPALITY	ARRIVAL Airport Municipality Derived from Code
ARRIVAL_AIRPORT_REGION	ARRIVAL Airport Region Derived from Code
DEPARTURE_AIRPORT_COUNTRY_CODE	DEPARTURE Airport Country code ISO2
DEPARTURE_AIRPORT_IATA	DEPARTURE Airport 3 character code
DEPARTURE_AIRPORT_MUNICIPALITY	DEPARTURE Airport Municipality Derived from Code.

 Target

 Features

Lab Tips

- Watson Studio url: dataplatform.ibm.com
- Labs are in www.github.com/bleonardb3/DS_POT_04-30-2020 repository.
- Instructions for each Lab are in the [README](#) file in the respective Lab folder.
- Cloud development enables making frequent improvements in the user interface. We reviewed the lab instructions and made screen updates so they should be pretty faithful to the user interface. Small differences may occur but shouldn't get in the way of successfully completing the labs.
- Do not use Internet Explorer or Edge as the browser. For Mac users do not use Safari.
- All of the Labs should be done in the project that you created in Lab-1

Github Repository

[bleonardb3 / DS_POT_04-30-2020](#)

Unwatch ▼

1

★ Star

0

Fork

0

<> Code

Issues 0

Pull requests 0

Actions

Projects 0

Wiki

Security 0

Insights

Settings

No description, website, or topics provided.

Edit

[Manage topics](#)

51 commits

1 branch

0 packages

0 releases

1 contributor

Branch: master ▼

New pull request

Create new file

Upload files

Find file

Clone or download ▼



bleonardb3 Update README.md

Latest commit 007d63c 5 hours ago

Lab-1	Update README.md	6 hours ago
Lab-2	Add files via upload	23 hours ago
Lab-3	Update README.md	3 days ago
Lab-4	Add files via upload	3 days ago
Lab-5	Update README.md	2 days ago
Lab-6	Add files via upload	2 days ago
Lab-7	Update README.md	23 hours ago
Lab-8-Appendix	Add files via upload	5 hours ago

Github Repository

Readme

1. [Lab-1](#) - This lab will set up the Watson Studio environment for subsequent labs and introduce you to the Project and Gallery features of Watson Studio
2. [Lab-2](#) - This lab will introduce you to the features of IBM's Watson Knowledge Catalog. Watson Knowledge Catalog is a secure enterprise catalog to discover, catalog and govern your data and modeling assets with greater efficiency.
3. [Lab-3](#) - This lab will introduce the Data Refinery. Data Refinery is a self-service data preparation tool for data scientists, data engineers, and business analysts. Data Refinery provides profiling, visualization, and a robust set of transforms to prepare data for analytics purposes. We will continue to use the 3 Trafficking data sets in this lab to demonstrate data profiling, data visualization, and data preparation capabilities of the Data Refinery tool. Note the datasets use simulated data.
4. [Lab-4](#) - In this lab, you will use the Watson SPSS Modeler capability to explore, prepare, and model the trafficking data. The SPSS Modeler is a drag and drop capability to build machine learning pipelines.
5. [Lab-5](#) - In this lab, you will use SparkML in Watson Studio to run simulated travel data through a machine learning algorithm, automatically tune the algorithm, and load the data into a DB2 Warehouse database.
6. [Lab-6](#) - This lab consists of two parts. The first part will demonstrate the new and exciting AutoAI capability to build and deploy an optimized model based on the trafficking data sets. The second part will deploy an application using the IBM Cloud DevOps toolchain that will invoke the deployed model to predict the human trafficking risk.
7. [Lab-7](#) - This lab will feature Watson OpenScale. IBM Watson OpenScale is an open platform that helps remove barriers to enterprise-scale AI.
8. [Lab-8-Appendix](#) - This lab will fulfill the prerequisites for Labs 3,4, and 5, if Lab-2 is not completed successfully.

Github Repository

Lab-1 Readme

Lab-1: Setup Environment

Introduction:

This lab will set up the Watson Studio environment for subsequent labs and introduce you to the Project and Gallery features of Watson Studio. Watson Studio is an integrated platform of tools, services, data, and meta-data to help companies and agencies accelerate their shift to be data driven organizations. The platform enables data professionals such as data scientists, data engineers, business analysts, and application developers collaboratively work with data to build, train, deploy machine learning and deep learning models at scale to infuse AI into business to drive innovation. Watson Studio is designed to support the development and deployment of data and analytics assets for the enterprise.

Objectives:

Upon completing the lab, you will:

1. Create a project
2. Create an object storage instance and associate it with the project
3. Create a Watson Machine Learning service instance and associate it with the project
4. Add a collaborator to the project
5. Research topics by searching the Gallery

Instructions:

Step 1. Please click on the link below to download the instructions to your machine.

[Instructions.](#)

Browser Tabs



Lab-1: Set up Environment

Introduction:

This lab will set up the Watson Studio environment for subsequent labs and introduce you to the Project and Community features of Watson Studio.

Objectives:

Upon completing this lab, you will know how to:

- Create a project
- Create an object storage instance and associate it with the project
- Create a Watson Machine Learning service instance and associate it with the project
- Add a collaborator to the project
- Research topics by searching the Gallery

Lab-2: Introduction to Watson Knowledge Catalog

Introduction:

This lab will introduce you to the features of IBM's Watson Knowledge Catalog. Watson Knowledge Catalog is a secure enterprise catalog to discover, catalog and govern your data and modeling assets with greater efficiency.

Objectives:

The goal of the lab is to gain familiarity with the features of the Watson Knowledge Catalog. Upon completing the lab, you will know how to:

- Create a governed catalog
- Add a member to the catalog
- Add Data Assets to the catalog
- Search the catalog
- Edit/Review/Profile a Data Asset
- Demonstrate access control features
- Create and enforce policy
- Push the Data Assets to a project.

End-to-End Data Science using
Watson Studio

Lunch Break. We will return for
review at 12:30 pm.

Lab-3: Introduction to the Data Refinery

Introduction:

In this lab, you will use the Watson Studio Data Refinery to profile data, visualize data, and prepare data for modeling.

Objectives:

Upon completing the lab, you will know how to:

- Profile the data
- Visualize the data to gain a better understanding
- Prepare the data for modeling
- Run the sequence of data preparation operations on the entire data set.

Lab-4: SPSS Modeler

Introduction:

In this lab, you will use the Watson Studio SPSS Modeler capability to explore, prepare, and model trafficking data. The SPSS Modeler is a drag and drop capability to build machine learning pipelines.

Objectives:

Upon completing the lab, you will:

- Become familiar with the Watson Studio SPSS Modeler capability
- Profile the data set
- Explore the data set with visualizations
- Transform the data
- Train/Evaluate a machine learning mode.

End-to-End Data Science using Watson Studio

We will return for lecture at
2:00 pm. Please work on labs 3 and 4

Categories of Machine Learning

Supervised learning


- The program is “trained” on a pre-defined set of “training examples”, which then facilitate its ability to reach an accurate conclusion when given new data
- The algorithm is presented with example inputs and their outcomes (labels)
- The goal is to learn a general rule that maps inputs to outputs

Unsupervised learning


- No labels are given to the learning algorithm, leaving it on its own to find structure (patterns and relationships) in its input

Supervised vs. Unsupervised Learning

Supervised Learning




Unsupervised Learning




Supervised vs. Unsupervised Learning

Supervised Learning



Unsupervised Learning



Preprocessing: Matrix for Machine Learning

Known as:


- Attributes
- Features
- Predictor variables
- Explanatory variables

Scale variables:

- Continuous variables, which can be measured on an interval scale or ratio scale
- 'Weight', 'Temperature', 'Salary', etc...

Categorical variables:

- Data with a limited number of distinct values or categories (nominal or ordinal)
- 'Hair color', 'Gender', 'Grape varieties', etc...



Known as:

- Label
 - Target variable
 - Dependent variable
- Scale or Categorical

Training, testing, & validation sets

During the model development process, supervised learning techniques employ **training** and **testing** sets and sometimes a **validation** set.

- Historical data with known outcome
- Data is randomly split into training, testing, and/or validation sets (mutually exclusive records)

Why?

- Training set
 - Build the model
 - Tune the parameters
- Validation set
 - Assess model quality during training/tuning process
 - Avoid overfitting the model to the training set
- Test set
 - Estimate accuracy or error rate of model after tuning
 - Used to compare multiple models

K-Fold Cross Validation

- **Instead of using a separate validation set**
- **Shuffle Training Samples and sub-divide into “K” folds (groups)**
- **Train “K” models using K-1 folds as training data and 1 Fold as Test Data**
- **For example, K=4**
 - Model 1 Train on 1,2,3 Test on 4 – calculate and store E1 (Error)
 - Model 2 Train on 2,3,4 Test on 1 – E2
 - Model 3 Train on 3,4,1 Test on 2 - E3
 - Model 4 Train on 4,1,2 Test on 3 - E4
 - $E = (E1+E2+E3+E4)/4$
- **A common value for K is 10**

Spark and Spark ML

Spark – why should I use it?

- Spark is a highly scalable runtime environment for analytics
- Provides the runtime engine and API
- Supports multiple languages: Python (PySpark), R (SparkR) and Scala

If you want to take advantage of Spark scalability and performance, you have to use Spark APIs

- Example (Python): Spark data frame vs. Pandas, Spark algorithms vs. scikit-learn
- It's possible to “mix and match” Spark and non-Spark code in a single notebook: the runtime environment will switch automatically
 - For example, use Python API for data understanding and SparkML for modeling

Spark Machine Learning API: <https://spark.apache.org/docs/latest/ml-guide.html>

Supported versions of Spark:

<https://www.ibm.com/software/reports/compatibility/clarity/prereqsForProduct.html>

Lab-5: Flow

Read in data from Cataloged Assets

- Join trafficking, job categories, occupations data

Identify Labels

- Label the data (“VETTING_LEVEL”)
- Select features

Feature Engineering (Transformation)

- StringIndexer (occupation, country, gender, birth year variables)
- VectorAssembler
- Normalizer

Define Model and Setup Pipeline

- Naïve Bayes
- Random Forest

Train the Model

- Split input data into Training (70%) and Test (30%) DataFrames
- Cache the resulting DataFrames
- Fit the Pipeline to the Training data set



Lab-5: Flow (continued)

Evaluate the resulting predictions

- Area under the ROC curve

Tune the model (hyperparameters)

- Build Parameter Grid
- Cross-evaluate to find the best model

Score the unvetted records

- Use Best Model to Score unvetted records (VETTING LEVEL == 100)
- Write results into the Database

Save the model in the Model Repository

- Model properties can be saved as well (e.g Area under the ROC curve)

Lab-5: Machine Learning using SparkML

Introduction:

In this lab, you will use SparkML in Watson Studio to run generated travel data through a machine learning algorithm, automatically tune the algorithm, and load the prediction results into a DB2 on Cloud database.

Objectives:

Upon completing the lab, you will know how to use a Jupyter Notebook to:

- Connect to a cataloged assets to read in data used for machine learning.
- Select the target and features
- Transform data
- Declare a machine learning model.
- Setup up the data transform and modeling pipeline
- Train the model.
- Evaluate the model.
- Automatically tune the model.
- Score data and load into a new DB2 table.
- Save the trained model

Lab-6: AutoAI + DevOps

Introduction:

This lab consists of two parts. The first part will demonstrate the exciting AutoAI capability to build and deploy an optimized model based on the trafficking data set. The second part will deploy a web application using the IBM Cloud DevOps toolchain that will invoke the deployed model to predict the trafficking risk.

Objectives:

Upon completing the lab, you will:

- Become familiar with the AutoAI feature of Watson Studio.
 - Train/Evaluate a machine learning model
 - Deploy a machine learning model.
-
- Deploy a Python Flask web application that we will configure to "call" the deployed machine learning model.
 - Configure the application to connect to the machine learning service.
 - Update the code in the application to specify the endpoint of the deployed model, and use DevOps to build and re-deploy the application.
 - Run the application to demonstrate the use of the deployed machine learning model to score the trafficking risk of a passenger.

End-to-End Data Science using Watson
Studio

We will return for lecture at
3:30 pm. Please work on labs 5 and 6


Trusted AI



**Is it
accurate?**



**Is it
fair?**



**Is it easy to
understand
?**



**Did
anyone
tamper
with it?**

Watson OpenScale



- **Trust and Transparency**

- Intelligently delivers bias mitigation help
- Provides traceability & auditability of AI predictions made in production applications
- Tracks AI accuracy in applications
- Explains an outcome in business terms
- Provides drift detection


- **Automation**

- Automatically detects and mitigates bias in model output, without affecting currently deployed model or outcomes

- **Open by Design**

- Monitor models deployed on third party model server engines
- Deploy behind enterprise firewall or on IaaS provider.

Model build / train frameworks



Model serving environments



Operations dashboard

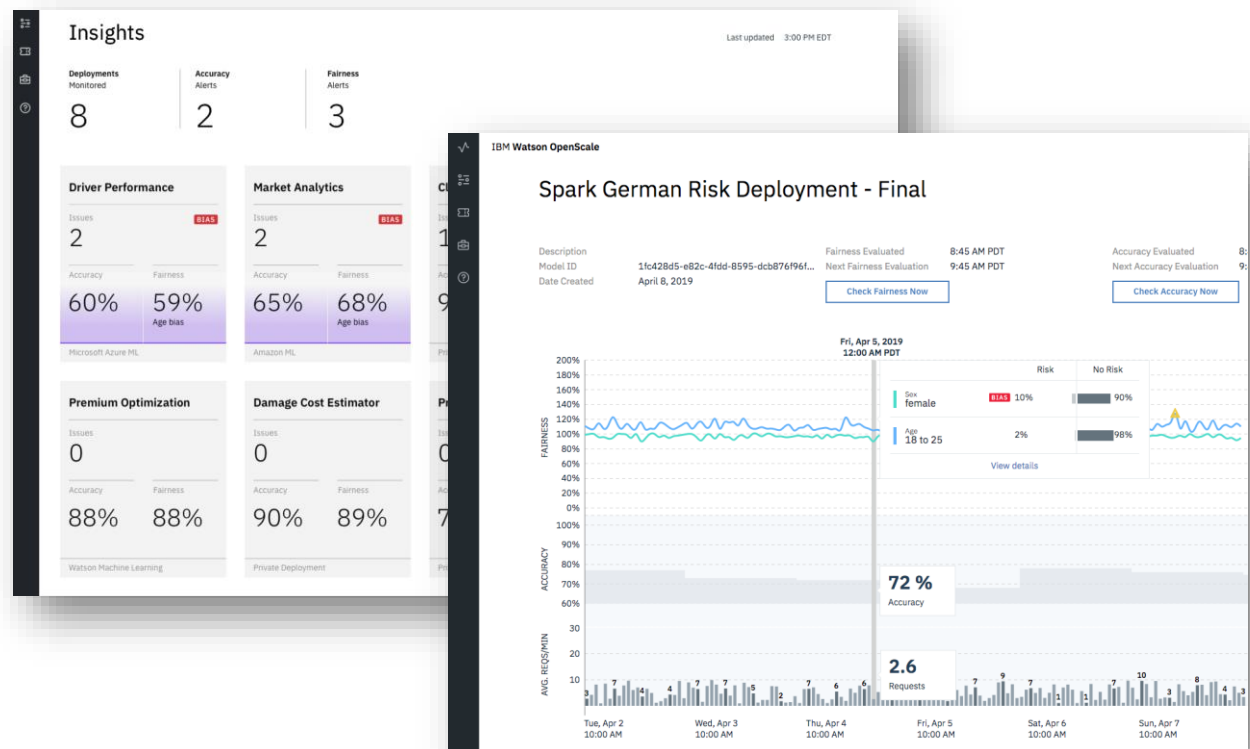


Description:

Monitor deployed models in a single dashboard that can be filtered by deployment making it easy to manage AI in apps

Value:

- Configure alerts or actions to be triggered when KPIs exceed threshold, ensuring model quality for improve business outcomes
- Measure model accuracy as it pertains to it's ability to deliver outcomes more accurate than knowledge workers
- Provides “continuous evolution” for your models



Model Fairness

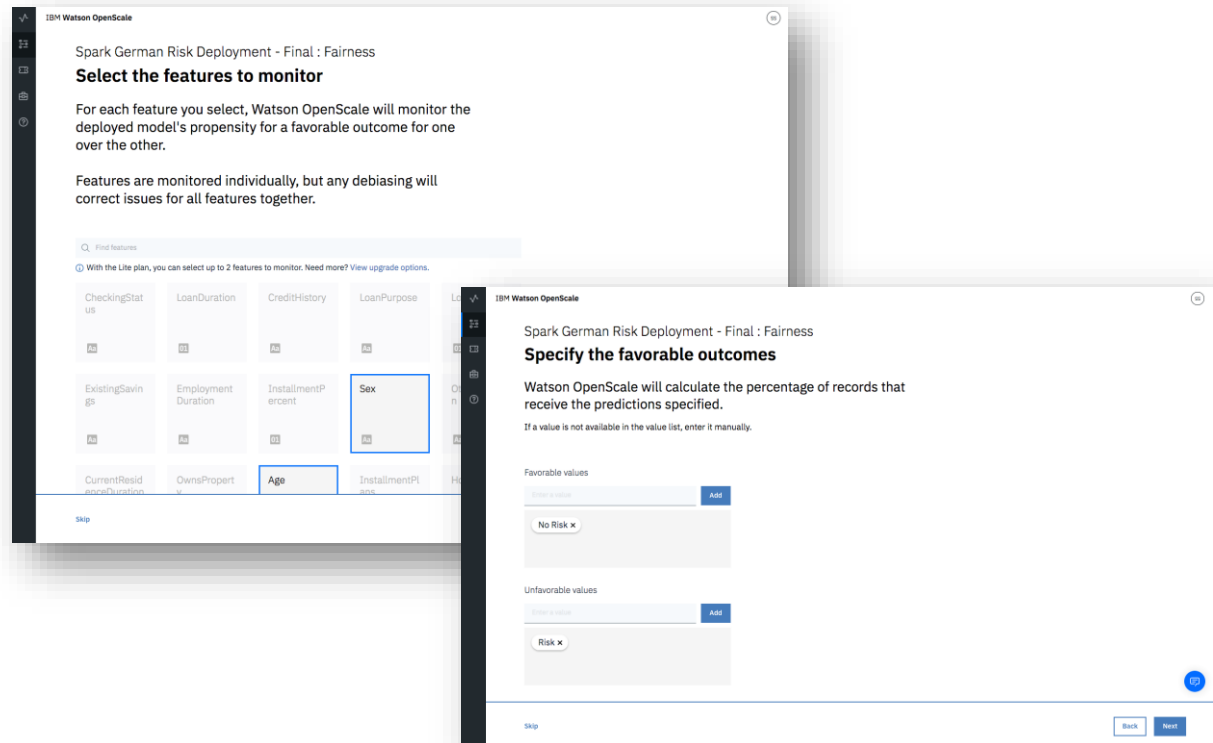


Description:

Production Models need to make fair decisions and *can not be biased* in their recommendations

How it works:

- Outcomes are selected as “favorable or unfavorable”
- “Favored Populations” and “protected populations” are selected where majority and minority groups are found
- A score is calculated based on the probability of favorable outcome for minority vs. probability of favorable outcome for majority



The image displays two overlapping screenshots of the IBM Watson OpenScale user interface. The background screenshot shows the 'Select the features to monitor' step for a 'Spark German Risk Deployment - Final : Fairness' model. It lists various features like 'CheckingStatus', 'LoanDuration', 'CreditHistory', 'LoanPurpose', 'ExistingSavings', 'EmploymentDuration', 'InstallmentPercentage', 'Sex', 'CurrentResidenceDuration', 'OwnsProperty', 'Age', and 'InstallmentPeriod'. The 'Sex' feature is highlighted with a blue border. The foreground screenshot shows the 'Specify the favorable outcomes' step for the same model. It prompts the user to specify favorable and unfavorable values. Under 'Favorable values', 'No Risk' is selected. Under 'Unfavorable values', 'Risk' is selected. Both screenshots include a 'Skip' button at the bottom left and 'Back' and 'Next' buttons at the bottom right.

Explainability



Description:

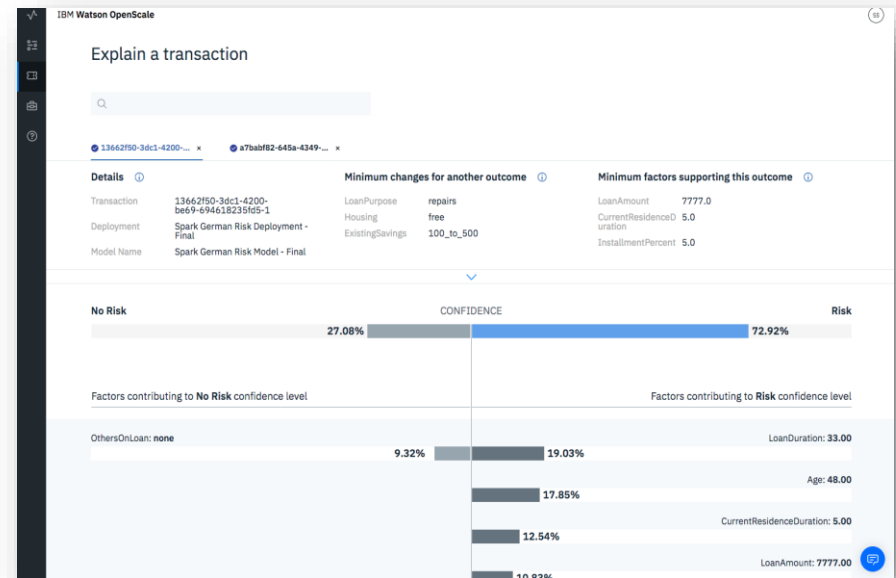
Allows you to understand which feature values of a model that are most influencing a prediction for a specific transaction

Example:

A loan is not approved by a model prediction - explainability will tell you why

How it works:

- Perturbation analysis on thousands of variations
- Risk model is created for two variations:
 - **LIME (local) Explanation:** set of features which played a positive or negative role in the prediction - also identifies the feature weights which helps to identify the most or least important features
 - **Contrastive Explanation:** Explains the behavior of the model in the vicinity of the data point whose explanation is being generated – assumption: the most common value is the least interesting from an explanation point of view



Lab 7: Watson OpenScale

Introduction:

IBM Watson OpenScale is an open platform that helps remove barriers to enterprise-scale AI. In this lab you will configure Watson OpenScale to monitor quality, fairness, and drift and to provide the factors that explain a deployed model's classification.

Objectives:

Upon completing the lab, you will

- Import and Deploy a machine learning model
- Provision a Watson OpenScale service
- Configure Watson OpenScale for Payload Logging, Quality, Fairness, and Drift.
- Submit Feedback and View Quality Metrics
- Score Data and View Fairness Metrics
- Explain a Transaction.

Watson Studio supports the Data Science Lifecycle

Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.

